

D-MeZO-N: Decentralized Federated MeZO with Nesterov-Style Stabilization

Maxim Sukhatsky — Bauman MSTU (Kaluga branch) — loxtrepoi@gmail.com — github.com/Siesher/dmezo
Spring 2026

Abstract

We introduce **D-MeZO-N** — Decentralized Federated MeZO with Nesterov-style **stabilization** — a peer-to-peer federated zeroth-order optimizer for large language model fine-tuning with formal momentum stability analysis under bounded variance. Building on Malladi et al.’s MeZO (memory-efficient zeroth-order, NeurIPS 2023) we replace the single-machine setup with n clients connected by a doubly-stochastic mixing matrix W (Koloskova et al. 2020), where each client **independently** samples its own perturbation direction z_i and communicates only a single scalar (ρ_i) and integer seed per round per neighbour — eliminating gigabyte-scale gradient exchange and providing $1/n$ variance reduction across **both** data and direction noise components. To stabilise heavy-ball momentum under high ZO variance we introduce **adaptive ρ -clipping** (running 95-percentile $\times 1.3$) combined with **drift-reset** (zeroing velocity on eval-loss uptick detection) and linear β -decay $0.9 \rightarrow 0$ — the D-MeZO-N v2 (combo) recipe. **On Qwen3.5-4B-Base / MathLogicQA / 3 paired seeds D-MeZO-N v2 achieves final loss 1.293 ± 0.010 vs vanilla MeZO 1.368 ± 0.018 ($\Delta = -5.5\%$, 3/3 same direction, paired $t = 5.3$, $p < 0.05$)** — the first paper-scale multi-seed validated empirical improvement of decentralized federated MeZO over vanilla MeZO on a convergent reasoning task. Final accuracy shows a positive but statistically insignificant trend (+2.3pp mean; per-seed $\{-1, +8, 0\}$ pp; paired $t \approx 0.8$, $p \approx 0.5$; 100-example eval pool, $SE \approx 4.9$ pp) — we report it for transparency and do not claim it as a gain. Component recipes either fail multi-seed validation (v1 fixed clip $C=50$: 3/3 worse, +7.0% loss; drift-reset alone: 3/3 worse, +6.4% loss) or show seed-specific accuracy effects (adaptive clip alone: loss winner, acc varies). Combo achieves the **lowest loss-std among the variants that improve on vanilla** (combo 0.010 vs adaptive-clip-alone 0.021; vanilla itself 0.018) — an additional stability metric. (The drift-only variant has an even smaller loss-std of 0.0035, but this reflects that it is uniformly stuck at a worse loss, +6.4%, not the stability of a good solution.) We complement the empirical study with four formal convergence theorems — **Theorem 1** (convex + momentum + decentralized, ρ -clipping), **Theorem 2** (non-convex Polyak-Łojasiewicz, no momentum, $1/n$ federated speedup), **Theorem 3** (PL + heavy-ball + β -decay + ρ -clip; Lyapunov technique — to our knowledge the first stability guarantee for ZO heavy-ball in the PL regime) and **Theorem 4** (DP extension of T3 via dual-use ρ -clip as L2-sensitivity; per-round ($\epsilon=10$, $\delta=10^{-3}$)-DP at $\sim 6\%$ utility cost on Qwen3.5-0.8B / MathLogicQA). **Importantly, Theorem 3 establishes stability and linear convergence to a noise floor; the contraction rate $(1 - 3\eta\mu/2)$ does not improve on the plain SGD-under-PL rate $(1 - 2\eta\mu)$ (Karimi et al. 2016) — no asymptotic acceleration is claimed**, consistent with Bottou–Curtis–Nocedal 2018 (Thm 5.1). A $3\times$ empirical transient speedup (Day 8 R1b) is reported as observation; a formal transient acceleration analysis remains open. Full proofs are consolidated in docs/theory_rigorous.md.

1. Introduction

Memory-efficient zeroth-order (MeZO) optimization of large language models was introduced by Malladi et al. (2023) as a surprising result: fine-tuning a multi-billion-parameter LLM requires only forward passes, with a memory cost equal to inference. The key trick — replacing backpropagation with a two-point gradient estimator over a random direction reconstructable from a seed — drops the optimizer state from $O(d)$ (Adam moments) to $O(1)$ (a single scalar). For

federated learning this property is transformative: instead of streaming dense gradients between clients, MeZO clients exchange only (s, ρ) pairs.

This paper closes the gap with eight contributions:

- **C1** — First federated MeZO on a hybrid linear-attention LLM (Qwen3.5-4B-Base).
- **C2** — D-MeZO is robust to extreme partition heterogeneity: Dirichlet($\alpha = 0.5$) tax $\leq 13\%$ on the Day 5 grid.
- **C3** — Topology cost $\leq 8\%$ at $n = 4$ clients; counter-intuitively, ring(4) $\leq$ complete(4) on the ZO regime.
- **C4** — **D-MeZO-N v2 = combo (adaptive p-clip B1 + drift-reset B5)** on paper-scale (Qwen3.5-4B-Base / MathLogicQA / 3 paired seeds): final loss 1.293 ± 0.010 vs vanilla 1.368 ± 0.018 ($\Delta = -5.5\%$, 3/3 same direction, paired $t = 5.3$, $p < 0.05$). Accuracy is reported for transparency only — a positive but statistically insignificant trend (+2.3pp mean, per-seed $\Delta \in \{-1, +8, 0\}$ pp; paired $t \approx 0.8$, $p \approx 0.5$; SE ≈ 4.9 pp on the 100-example eval pool) — not claimed as a gain. Lowest loss-std among the variants that improve on vanilla (0.010 vs adaptive_clip alone 0.021; vanilla itself 0.018 — drift-only’s smaller 0.0035 reflects being uniformly stuck at a worse loss, not stability). **First paper-scale multi-seed validated empirical improvement of decentralized federated MeZO over vanilla MeZO (loss).**
- **C5** — Independent z_i per client (vs FedKSeed shared- z) provide $1/n$ variance reduction in **both** data and direction noise. Communication: 16 bytes/round/neighbor at `update_share` consensus mode.
- **C6** — **Theorem 3**: closed-form Lyapunov convergence for PL + heavy-ball + ρ -clip + β -decay at rate $(1 - 3\eta\mu/2)$ and neighbourhood $2G^2/(3\mu)$ — to our knowledge the **first stability guarantee for ZO heavy-ball (PL regime)**. **Importantly: the contraction rate $(1 - 3\eta\mu/2)$ does not improve on the plain SGD-under-PL rate $(1 - 2\eta\mu)$ (Karimi et al. 2016) — no asymptotic acceleration claimed** (consistent with Bottou–Curtis–Nocedal 2018, Thm 5.1); the transient empirical speedup remains an open question.
- **C7** — **Theorem 4**: DP extension of T3 via **dual-use p-clip** — same C used for momentum stability **simultaneously serves as L2-sensitivity** for Gaussian mechanism. Per-round ($\varepsilon = 10$, $\delta = 10^{-3}$)-DP at $\sim 6\%$ utility cost on Qwen3.5-0.8B / MathLogicQA. T-round composition acknowledged as explicit limitation; subsampling amplification (Abadi 2016) — future work.
- **C8** — Honest multi-seed validated negatives: v1 (fixed $C = 50$): 3/3 worse (+7.0% loss); drift-reset alone: 3/3 worse (+6.4%); look-ahead Nesterov diverges $7\times$ faster (R20 vs R140); K=3 multi-direction equal-compute loses; $\varepsilon(t)$ warmup schedules robustly lose across 16+ cells. Clean falsification evidence on scope of applicability.

2. Related Work

SPSA — foundational ancestor. Simultaneous Perturbation Stochastic Approximation (Spall 1992, IEEE TAC 37:332) is the two-point zeroth-order optimizer underlying all subsequent ZO literature: $\theta \leftarrow \theta - \eta \cdot ((L(\theta + \varepsilon\Delta) - L(\theta - \varepsilon\Delta)) / (2\varepsilon)) \cdot \Delta^{-1}$ with Bernoulli Δ . Convergence rate $O(1/\sqrt{T})$ (Sadegh-Spall 1998); 2-SPSA with Hessian estimation (Spall 2000); adaptive gain $\eta_t = a/(t + A)^\alpha$ with $\alpha \approx 0.602$ (Spall 1997). Classical SPSA does not scale to LLMs because the perturbation Δ must be explicitly stored ($O(d)$ memory) — this gap is closed by MeZO.

MeZO — SPSA for LLMs. Malladi et al. (2023, NeurIPS) introduced MeZO as SPSA with two key innovations: (i) **Gaussian** $z \sim \mathcal{N}(0, I)$ instead of Bernoulli (enables the $r(H)$ -trick below), (ii) **in-place seed-reconstructed** z — direction z is deterministically regenerated from a seed rather than materialized as a tensor, keeping memory at inference level. Theorem 3.1 proves a variance bound that uses the effective Hessian rank $r(H) := \text{tr}(H)/\|H\|_{op}$ instead of full dimension d , justifying applicability at LLM scale ($r(H) \ll d$ for overparametrized transformers).

ℓ_2 -smoothing analysis (Russian school). Parallel to SPSA, a Gaussian/uniform-smoothing tradition for convex ZO has been developed (Nesterov-Spokoiny 2017, FOCM). Gasnikov et al. (2023, “Randomized gradient-free methods in convex optimization”, *Encyclopedia of Optimization*, arXiv:2211.13566) systematise this framework: two-point estimator $\nabla f_\gamma(x, e) = d \cdot \frac{f(x+\gamma e) - f(x-\gamma e)}{2\gamma} \cdot e$ for $e \sim \text{Unif}(S_2^d)$, bias bound $f(x) \leq f_\gamma(x) \leq f(x) + \gamma M_2$, and variance bound $\mathbb{E} \|\nabla f_\gamma\|^2 \leq \sqrt{2} \min(q, \ln d) d^{2/q} M_2^2$. These bounds form the formal foundation of our convergence proofs T1-T4. The survey covers convex/strongly convex cases with accelerated rates; PL + heavy-ball (where D-MeZO-N operates) is not addressed — this is our delta.

D-MeZO-N — SPSA + MeZO + 4 new elements. Our method builds on the SPSA lineage (34 years of stochastic approximation theory) and adds: (1) federated wrapper with **independent per-client** z_i plus gossip mixing matrix W (vs FedKSeed-style shared seed pool); (2) heavy-ball scalar momentum with adaptive ρ -clip and drift-reset (vs SPSA’s adaptive gain without momentum); (3) closed-form Lyapunov convergence proof under PL — to our knowledge the first stability guarantee for ZO heavy-ball in the PL regime; (4) dual-use ρ -clip as L2-sensitivity for Gaussian-mechanism DP (a targeted literature search for “DP-SPSA” returns no results — open niche).

Decentralised SGD. Koloskova et al. (2020) provide a unified analysis for D-SGD with arbitrary mixing matrices W . Their Theorem 2 (convex) and Theorem 8 (PL) bound the convergence rate as a function of the spectral gap $\rho(W)$ and gradient heterogeneity ζ^2 ; the framework subsumes local-update methods (Stich 2019). In parallel, Beznosikov, Richtárik et al. (NeurIPS 2022, arXiv:2207.10792) develop decentralized SGD with communication compression for variational inequalities — a VI-counterpart to our 16-byte-per-round economy; Sadiev et al. (2022, *EURO J Comp Opt*) provide lower bounds for decentralized personalized FL that bound the optimality of our T1 from below.

Federated zeroth-order. FedKSeed (Qin et al., 2024 ICML), Ferret (Shu et al., 2024) and FedZeN (Maritan et al. 2024) all build on MeZO for FL, but are limited to (i) full-attention architectures and (ii) FedAvg-style central-server aggregation. DeComFL (Li et al. 2024, ICLR 2025) likewise compresses FL communication to $O(1)$ scalars per round via shared-seed zeroth-order steps and proves dimension-free rates under a low-effective-rank assumption, but remains server-orchestrated with one per-round seed shared across clients. D-MeZO-N differs on all three axes: peer-to-peer gossip over any doubly-stochastic W , **independent** per-client z_i (the source of the $1/n$ reduction in both noise components, §4.4), and a momentum-stability analysis (Theorem 3).

Heavy-ball under PL. Yang, Lin, Li (2016) give a unified Lyapunov analysis of heavy-ball SGD in convex and non-convex PL regimes. Karimi, Nutini, Schmidt (2016) establish the canonical linear-convergence-to-noise-floor framework for stochastic gradient methods under PL.

3. Method: D-MeZO-N

3.1 Setup

Let n clients each hold a local data shard D_i and a local copy of the model parameters $\theta_i \in \mathbb{R}^d$. The federation is described by a doubly-stochastic mixing matrix $W \in \mathbb{R}^{n \times n}$ with spectral gap

$$\rho(W) := \|W - \frac{1}{n} \mathbf{1}\mathbf{1}^\top\|_{op} \in [0, 1).$$

3.2 Algorithm

On round t , each client i performs a MeZO step using a fresh seed s_i^t , producing the projected gradient

$$\hat{g}_i^t = \frac{L(\theta_i^t + \epsilon z) - L(\theta_i^t - \epsilon z)}{2\epsilon}.$$

The full D-MeZO-N round combines ρ -clip, a heavy-ball (Polyak 1964) velocity update with momentum β_t , a parameter step, and a consensus mixing step:

$$\begin{aligned} v_i^{t+1} &= \beta_t v_i^t + \text{clip}(\hat{\rho}_i^t, \pm C) z_{s_i^t}, \\ \theta_i^{t+1/2} &= \theta_i^t - \eta v_i^{t+1}, \\ \theta_i^{t+1} &= \sum_j W_{ij} \theta_j^{t+1/2}. \end{aligned}$$

Figure 5. D-MeZO-N algorithm — decentralized peer-to-peer with consensus mixing




Figure 1. D-MeZO-N algorithm for $n = 4$ clients on a ring topology. Each client performs an independent local MeZO probe (seed s_i , scalar ρ_i), clips ρ_i , updates a local velocity buffer with the scheduled β_t , then participates in a doubly-stochastic consensus mixing step.

3.3 ρ -clipping

We bound the per-step contribution to v_i by symmetric clipping:

$$\text{clip}(x, \pm C) := \max(-C, \min(C, x)).$$

A fixed threshold $C = 50$ (v1 recipe) was selected empirically on the SST-2 ablations (§5.4). At paper scale this fixed value proved over-tight (§5.6); the **D-MeZO-N v2** recipe replaces it with an adaptive threshold $C_t = \alpha \cdot$

$\text{quantile}_{0.95}(\{|\hat{\rho}|\}_{\text{last 50 rounds}})$, $\alpha = 1.3$ (component B1), paired with a drift-reset that zeros the velocity buffer on eval-loss uptick (component B5). See §5.6.2 for the v2 derivation and headline result.

4. Theory

4.1 Assumptions

- (A1) Smoothness: each L_i is ℓ -smooth, $\|\nabla L_i(\theta) - \nabla L_i(\theta')\| \leq \ell \|\theta - \theta'\|$ (we reserve L for the loss and ℓ for the smoothness constant, following docs/theory_rigorous.md).
- (C2) Bounded gradient diversity: $\frac{1}{n} \sum_i \|\nabla L_i(\theta) - \nabla L(\theta)\|^2 \leq \zeta^2$.

- **(C3)** Bounded stochastic noise: $\mathbb{E}_\xi \|\nabla \ell(\theta; \xi) - \nabla L_i(\theta)\|^2 \leq \sigma_b^2$.
- **(C5)** Effective Hessian rank: $r(H) := \text{tr}(H)/\|H\|_{op} \ll d$.
- **(A2 / PL, used in Theorem 2 only):** $\|\nabla L(\theta)\|^2 \geq 2\mu(L(\theta) - L^*) \quad \forall \theta \in \mathbb{R}^d$.

4.2 Lemmas

Lemma 1 (Malladi ZO variance, H -weighted form). *Under (A1)+(C5), for any PSD matrix M the MeZO estimator $\hat{g} = \hat{\rho} z$ satisfies the quadratic-form bound (docs/theory_rigorous.md Lemma 1b'):*

$$\mathbb{E}_z [\hat{g}^\top M \hat{g}] \leq \|M\|_{op} \|\nabla L\|^2 (r(M) + 2) + O(\epsilon^2).$$

Instantiated at the local Hessian $M = H$ (so $\|H\|_{op} \leq \ell$, $r(M) = r(H)$), this gives the form used by the descent analysis:

$$\mathbb{E}_z [\hat{g}^\top H \hat{g}] \leq \ell (r(H) + 2) \|\nabla L\|^2 + O(\epsilon^2 \ell_3(d + 2)),$$

*where the bias term inherits the $O(\epsilon^2 \ell_3(d + 2))$ scale of Lemma 1a (docs/theory_rigorous.md Lemma 1a/1b'). Note that the **raw Euclidean second moment scales with the full dimension**, $\mathbb{E}_z \|\hat{g}\|^2 = (d + 2) \|\nabla L\|^2 + O(\epsilon^2)$ (docs/theory_rigorous.md Lemma 1b); the effective-rank reduction $r(H) \ll d$ applies only inside H -weighted quadratic forms, which is precisely where the descent analysis uses it.

Lemma 2 (ρ -clipping bias-variance). *Let $\tilde{\rho} = \text{clip}(\hat{\rho}, \pm C)$. Then*

$$\mathbb{E} \|\tilde{\rho} \cdot z\|^2 \leq \min(\mathbb{E} \|\hat{\rho} \cdot z\|^2, C^2 d).$$

Lemma 3 (consensus error). *For D-MeZO-N with mixing matrix W ($\rho := \rho_W < 1$) and momentum schedule $\beta_t \in [0, \beta_0]$ ($\bar{\beta} := \beta_0/2$) (docs/theory_rigorous.md Lemma 4):*

$$\frac{1}{n} \sum_i \|\theta_i^{t+1} - \bar{\theta}_{t+1}\|^2 \leq \frac{\rho^2}{(1 - \rho)^2} \cdot \frac{\eta^2 (G^2 r(H) + \zeta^2)}{(1 - \bar{\beta})^2}.$$

Lemma 5 (PL descent; Karimi-Nutini-Schmidt 2016). *Under (A1)+(A2)+(C2)+(C3) for $\eta \leq 1/(2\ell)$:*

$$\mathbb{E}[L(\theta_{t+1}) - L^*] \leq (1 - \eta\mu) \mathbb{E}[L(\theta_t) - L^*] + \frac{\eta^2 \ell \sigma_b^2}{2} + \frac{\eta \zeta^2}{\mu}.$$

(Stated in the paper's notation: σ_b^2 is the stochastic-noise bound (C3), ζ^2 the gradient-diversity bound (C2). Theorem 2 invokes this lemma at the tighter step size $\eta \leq 1/(4\ell)$, whose slack absorbs the ZO variance factor of Lemma 1 and yields the $(1 - \eta\mu/2)$ rate.)

4.3 Theorem 1 — convex case with momentum

Theorem 1 (D-MeZO-N convergence, convex case). *Assume (A1), (C2), (C3), (C5) with each L_i convex. The D-MeZO-N iterate satisfies:*

$$\mathbb{E}[L(\bar{\theta}_T) - L^*] \leq \tilde{O}\left(\sqrt{\frac{\ell \cdot r(H) \cdot D_0}{nT}}\right) + \tilde{O}\left(\frac{\rho^2}{(1 - \rho)^2} \cdot \frac{C^2 r(H)}{(1 - \bar{\beta})^2 T}\right) + O(\epsilon^2 \ell^2 r(H)),$$

where $D_0 = \|\bar{\theta}_0 - \theta^*\|^2$ and $\tilde{O}(\cdot)$ hides constants and logarithmic factors; the consensus-penalty term carries the mixing amplifier $\rho^2/(1 - \rho)^2$ (from `docs/theory_rigorous.md` Lemma 4 / line 361) and the momentum amplifier $1/(1 - \bar{\beta})^2$.

Proof sketch. Under bare convexity no PL constant μ exists, so the bound is *not* obtained by a geometric Lyapunov contraction. Instead we expand the distance to the optimum of the averaged iterate, $\|\bar{\theta}_{t+1} - \theta^*\|^2 = \|\bar{\theta}_t - \theta^*\|^2 - 2\eta \langle \bar{v}_{t+1}, \bar{\theta}_t - \theta^* \rangle + \eta^2 \|\bar{v}_{t+1}\|^2$, and apply the convexity inequality $\langle \nabla L(\bar{\theta}_t), \bar{\theta}_t - \theta^* \rangle \geq L(\bar{\theta}_t) - L^*$ to the gradient part of $\mathbb{E}\bar{v}_{t+1}$ (the residual is a bounded consensus-drift term ξ_t with $\|\xi_t\| \leq \ell\sqrt{\Pi_t}$). Summing over t , dividing by $2\eta T$, and applying Jensen to the Polyak-Ruppert average $\bar{\theta}_T = \frac{1}{T} \sum_t \bar{\theta}_t$ yields $L(\bar{\theta}_T) - L^* \leq D_0/(2\eta T) + (\text{drift term}) + (\text{variance term})$ with $D_0 = \|\bar{\theta}_0 - \theta^*\|^2$. The variance term is the $1/n$ -reduced per-client ZO floor (Lemma 1, H -weighted); bounding the drift via the consensus Lemma (`docs/theory_rigorous.md` Lemma 4) and optimizing the step size at $\eta^* = \sqrt{D_0 n / (C^2 r(H) \ell T)}$ produces the displayed $\sqrt{D_0 C^2 r(H) \ell / (nT)}$ rate. The full proof is in `docs/theory_rigorous.md` §4. ■

4.4 Theorem 2 — non-convex PL case (no momentum)

Theorem 2 (D-MeZO convergence, non-convex PL, $\beta = 0$). Under (A1)+(A2/PL)+(C2)+(C3)+(C5) and step-size $\eta \leq 1/(4\ell)$, for the complete-graph case ($\rho = 0$):

$$\mathbb{E}[L(\bar{\theta}_T) - L^*] \leq \left(1 - \frac{\eta\mu}{2}\right)^T \Delta_0 + \tilde{O}\left(\frac{\eta \ell C^2 r(H)}{\mu n}\right) + O\left(\frac{\epsilon^2 \ell^2 r(H)}{\mu}\right).$$

We prove Theorem 2 for the complete-graph case ($\rho = 0$; `docs/theory_rigorous.md` §2). For general doubly-stochastic mixing ($\rho > 0$) the analysis acquires an additional consensus-penalty term

$$\tilde{O}\left(\frac{\eta^2 \rho^2 \ell^2 C^2 r(H)}{\mu(1 - \rho)^2}\right),$$

which follows from the Theorem 1 decentralized machinery (the $\rho^2/(1 - \rho)^2$ mixing amplifier) and is stated here as a conjectured extension rather than a proved result.

4.5 Theorem 3 — PL case with heavy-ball momentum and β -decay

Theorem 3 (D-MeZO-N convergence under PL + momentum). Assume (A1)+(A2/PL) and a ρ -clipped second-moment bound (A4) $\mathbb{E}\|\tilde{\rho} z\|^2 \leq G^2$ for all iterations (under clipping at C , $G^2 \leq C^2 r(H) \ell$ in the Malladi-refined form — `docs/theory_rigorous.md` (A4)). For $\eta \leq (1 - \beta_0^2)/(8\ell)$ and any schedule $\beta_t \in [0, \beta_0]$ (including constant β_0 and linear decay $\beta_0 \rightarrow 0$), the heavy-ball MeZO iterate satisfies

$$\mathbb{E}[V_T] \leq \left(1 - \frac{3\eta\mu}{2}\right)^T V_0 + \frac{2G^2}{3\mu},$$

where $V_t := (L(\theta_t) - L^*) + (\eta/2)\|v_t\|^2$ is a Lyapunov function combining potential and kinetic energy. The loss component obeys the same bound:

$$\mathbb{E}[L(\theta_T) - L^*] \leq \left(1 - \frac{3\eta\mu}{2}\right)^T V_0 + \frac{2G^2}{3\mu}.$$

As stated, Theorem 3 covers the centralized / complete-graph iterate; the decentralized extension combining $\rho_W < 1$ mixing with heavy-ball under PL remains open (§7; `docs/theory_rigorous.md` §6).

Proof sketch. From the extended (A1)-smoothness in both components of V_t we obtain the descent inequality $V_{t+1} \leq V_t - \frac{3\eta}{4} \|\nabla L_t\|^2 - \frac{3\eta(1-\beta_t^2)}{8} \|v_t\|^2 + \eta G^2$ (the cross-terms $\pm \eta \beta_t \langle \nabla L_t, v_t \rangle$ from the potential and kinetic components cancel exactly — this is why the kinetic multiplier in V_t is $\eta/2$). The kinetic term is non-positive and contracts V_t at rate $3(1 - \beta_t^2)/4$, while PL converts the potential term, $-\frac{3\eta}{4} \|\nabla L\|^2 \leq -\frac{3\eta\mu}{2} (L - L^*)$. Taking the minimum of the two contraction rates and choosing $\eta = (1 - \beta_0^2)/(8\ell)$ (so the potential rate dominates for the admissible μ range) yields a uniform rate $3\eta\mu/2$ independent of the β_t schedule. Geometric summation $\sum_{k=0}^{\infty} (1 - 3\eta\mu/2)^k \eta G^2 = 2G^2/(3\mu)$ gives the neighbourhood. The full proof is in `docs/theory_rigorous.md` §3. ■

Corollaries.

- **Linear convergence** to the $2G^2/(3\mu)$ -neighbourhood. The contraction rate $(1 - 3\eta\mu/2)$ **does not improve on the plain SGD-under-PL rate** $(1 - 2\eta\mu)$ (Karimi et al. 2016) — to our knowledge this is the first such stability guarantee in the ZO-LLM setting, but it establishes that momentum can be made *safe* (non-divergent) in the high-variance ZO regime, not that it is *faster*. This is consistent with Bottou–Curtis–Nocedal (2018, Thm 5.1): heavy-ball gives no asymptotic acceleration under stochastic gradients.
- **Rescue mechanism.** Under ρ -clipping (A4), $G^2 \leq C^2 r(H) \ell$, so the neighbourhood is bounded; without clipping G^2 is unbounded and the iterate sequence diverges (§5.5 confirms this empirically).
- **β -decay.** The kinetic contraction term $3(1 - \beta_t^2)/4$ grows as $\beta_t \rightarrow 0$, removing the late-stage drift observed for constant β (R1b, §5.4 regions C/D).

4.6 Theorem 4 — DP extension of Theorem 3

We extend D-MeZO-N with isotropic Gaussian noise $\xi_t \sim \mathcal{N}(0, \sigma^2)$ added to the clipped projected gradient $\tilde{\rho}_t = \text{clip}(\hat{\rho}_t, \pm C)$, turning each MeZO step into one application of the Gaussian mechanism (Dwork-Roth 2014).

Theorem 4a (convergence under DP-noise). *Under the assumptions of Theorem 3 with DP-noise variance σ^2 , the Lyapunov iterate satisfies*

$$\mathbb{E}[V_T] \leq \left(1 - \frac{3\eta\mu}{2}\right)^T V_0 + \frac{2(C^2 r(H) + \sigma^2 d) \ell}{3\mu}.$$

*The two noise contributions scale with different dimensions (docs/theory_rigorous.md Lemma 8): the clip-noise $\tilde{\rho}z$ is aligned with the gradient and contributes $C^2 r(H)\ell$ (effective rank), whereas the DP-noise ξz is isotropic in parameter space (no alignment with ∇L) and contributes $\sigma^2 d\ell$ — scaling with the **full** dimension $d \gg r(H)$, i.e. DP-noise breaks the Malladi $r(H)$ -substitution trick. With this grouping, Theorem 4a reduces exactly to Theorem 3 (neighbourhood $2G^2/(3\mu)$ with $G^2 = C^2 r(H)\ell$) at $\sigma = 0$.*

Theorem 4b (per-round privacy). *One MeZO step with clip C and Gaussian noise σ provides (ε_1, δ) -DP per round with*

$$\varepsilon_1 = \frac{C\sqrt{2\ln(1.25/\delta)}}{\sigma}.$$

Proof: $\text{clip}(\hat{\rho}, \pm C)$ has L2-sensitivity $\Delta = C$; the Gaussian mechanism with $\sigma \geq \Delta\sqrt{2\ln(1.25/\delta)}/\varepsilon$ yields (ε, δ) -DP (Dwork-Roth 2014). The dual-use of C — momentum-stabilization threshold and L2-sensitivity bound simultaneously — means no extra per-sample gradient clipping is needed (cf. DP-SGD, Abadi et al. 2016). ■

Composition caveat (honest). The per-round ε_1 above does **not** compose for free over T training rounds:

- **Basic composition** (Dwork-Roth Thm 3.16): $\varepsilon_T = T\varepsilon_1$. For $T = 200$ and $\varepsilon_1 = 10$ this gives $\varepsilon_T \approx 2000$ — vacuous.
- **Advanced / RDP composition** (Dwork-Rothblum-Vadhan 2010; Mironov 2017; Abadi 2016): the Gaussian mechanism with L2-sensitivity C and noise σ is $(\alpha, \frac{\alpha C^2}{2\sigma^2})$ -RDP per round (note the sensitivity factor C^2 — the bare $\alpha/(2\sigma^2)$ coefficient holds only for unit-sensitivity), composing to $(\alpha, \frac{T\alpha C^2}{2\sigma^2})$ -RDP. These savings are meaningful **only in the $\varepsilon_1 \ll 1$ regime**: at $\varepsilon_1 = 10$ the advanced-composition term $T\varepsilon_1(e^{\varepsilon_1} - 1)$ explodes and basic composition governs.

We therefore report **per-round ε** (the standard convention for one-shot federated fine-tuning) and treat T -round composition as an explicit limitation. Achieving a useful training-level budget therefore requires per-round $\varepsilon_1 < 1$ (i.e. larger σ), which is the regime a full RDP-accountant treatment must target (future work), together with mini-batch subsampling amplification (Abadi 2016) — which our current full-batch-per-round setup does not exploit. The full statements and proofs are in `docs/theory_rigorous.md` §6.5.

5. Experiments

5.1 Setup and sanity checks (cross-architecture)

All experiments share the Princeton-default hyperparameters ($\eta = 3 \cdot 10^{-7}$, $\epsilon = 10^{-3}$, bf16), $n = 4$ clients with full-state (`weight_avg`) consensus, and MLflow tracking; compute is Colab Blackwell (RTX PRO 6000, 96 GB) plus a local RTX 5070 Ti for small-model ablations. Day-1 sanity runs confirm MeZO fine-tuning on both architectures: Qwen3-4B (full attention) drops SST-2 eval loss by 88.1% in 1000 steps, and Qwen3.5-4B-Base (hybrid linear-attention V-L, §3.1, vision tower frozen) by 94.7% — to our knowledge the first MeZO result on a hybrid linear-attention LLM. A 2×2 cross-architecture $\times$ cross-task grid (Qwen3-4B / Qwen3.5-4B-Base $\times$ SST-2 / BoolQ) passes in all four cells, closing the architecture-generality check (C1).

5.2 Federated 2×2 grid: topology × partition (Day 5)

Figure 1. Day 5 grid — federated D-MeZO trajectories per cell
(Qwen3.5-4B-Base / SST-2 / 4 clients / 1000 rounds; 2 seeds + centralized reference)




Figure 2. Per-cell trajectories of the Day 5 federated grid. All federated configurations consistently descend below the centralized baseline.

Config	Final eval (mean ± range/2)	Accuracy (mean %)	vs. centralized 0.1762
complete + IID	0.1348 ± 0.0051	96.56%	−23.5%
complete + Dir($\alpha = 0.5$)	0.1507 ± 0.0089	95.00%	−14.5%
ring + IID	0.1271 ± 0.0014	97.81% ★ best	−27.9%
ring + Dir($\alpha = 0.5$)	0.1402 ± 0.0029	95.63%	−20.4%
centralized (reference)	0.1762 (n=1)	95.63%	—
R1d (D-MeZO-N) on worst cell	0.1291 (single seed)	95.63%	−26.7%

5.3 Why does federated beat centralized?

Figure 3. Federated D-MeZO consistently beats centralized MeZO on Qwen3.5-4B-Base / SST-2




Figure 3. (a) Final eval loss of each federated configuration vs. the centralised MeZO baseline. (b) Final accuracy comparison.

The empirical ratio $0.1271/0.1762 = 0.722$ is **directionally consistent** with Theorem 2’s variance-floor reduction $\propto 1/n$ — when n clients each perform an independent MeZO probe with their own seed s_i and direction z_{s_i} , the consensus-averaging step amounts to an unbiased average of n independent unit-direction probes, reducing the variance floor by factor n . We avoid claiming “ $0.722 \approx 1/\sqrt{4}$ matches Theorem 1 rate”: the theoretical $1/\sqrt{nT}$ refers to the **convergence rate** in the convex case (on Polyak-Ruppert average), not the **ratio of final losses**, which is governed by the **noise-floor term** $\eta C^2 r(H) \ell / (\mu n)$ in Theorem 2. The two quantities have different mechanisms and should not be conflated. See [docs/theory_rigorous.md §5](#) for the full predictions-vs-empirics table.

5.4 Momentum ablations on the worst cell: the phase diagram (Days 6–8)

The momentum design of D-MeZO-N was developed through Day 6–8 ablations on the **worst federated cell** of the §5.2 grid (ring + Dir(0.5), Qwen3.5-4B-Base / SST-2, seed 42), which map a clean phase diagram (Figure 4). **Region A:** naive heavy-ball at constant $\beta = 0.9$ without clipping diverges catastrophically at R140 — ZO noise is amplified by the velocity buffer; a true look-ahead Nesterov variant diverges $7\times$ faster (R20), because both the probe location and the update direction depend on v_t (dual-channel noise, §6.2). **Region B:** loose clipping ($C = 200$) prevents the instant blow-up but slow-diverges from R500. **Region C:** tight clipping ($C = 50$) at constant $\beta = 0.9$ (run R1b) yields a genuine $3\times$ transient speedup — eval loss 0.119 at R300 vs the no-momentum control’s 0.138 at R1000 — but late-drifts to 0.225 as the velocity noise floor $\|v_\infty\|^2 \approx G^2/(1 - \beta^2)$ takes over. **Region D:** tight clipping plus linear β -decay $0.9 \rightarrow 0$ (run R1d) removes the drift: monotonic descent to 0.1291, 6.0% below the no-momentum control (0.1373) on the same cell (Figure 5). The C→D transition defines the **v1 recipe** (β -decay + fixed $C = 50$); Theorem 3’s corollaries predict both the region-C drift ($\|v_\infty\|^2 \approx 5G^2$ at $\beta = 0.9$) and its removal under β -decay ($\|v_T\|^2 \rightarrow G^2$). All region boundaries are single-seed observations; the v1 recipe is later falsified at paper scale (§5.6.1), motivating v2.

Figure 2. Nesterov-MeZO phase diagram on worst Day 5 cell (ring(4) + Dir($\alpha=0.5$))
All runs at seed=42 for bit-exact ablation; ρ -clip + β -decay required for stable acceleration




Figure 4. Phase diagram of Nesterov-MeZO variants on the worst federated cell.

Figure 4. D-MeZO-N (R1d) detailed trajectory — β -decay produces monotonic descent
Final 0.129 beats vanilla 0.137 by 6%; β linearly decays $0.9 \rightarrow 0$ over $T=1000$




Figure 5. D-MeZO-N (R1d) detailed trajectory vs. control.

5.5 Cross-task validation: HellaSwag (4-way commonsense reasoning)

We further test D-MeZO-N on **HellaSwag** (Zellers et al. 2019), 4-way commonsense reasoning — substantially harder than SST-2/BoolQ because endings are multi-token and require world knowledge inference, not lexical signals. Same setup: Qwen3-4B (full-attention transformer, bf16, Apache 2.0), $\eta = 3 \cdot 10^{-7}$, $\epsilon = 10^{-3}$, 1000 steps/rounds, 2000 train examples, 200 eval examples, seed=42.

Run	Init loss → Final loss	Δ loss	Init acc → Final acc	Δ acc	Verdict
Centralized vanilla MeZO	2.5691 → 2.7112	+5.5%	0.6625 → 0.6375	−2.50pp	DIVERGED
Federated D-MeZO-N v1 (4c complete IID, β -decay $0.9 \rightarrow 0$, ρ -clip $C = 50$)	2.5691 → 2.4959	−2.85%	0.6625 → 0.7000	+3.75pp	CONVERGED
Δ federated vs. centralized	−7.9% relative loss	—	+6.25pp absolute acc	—	—

Key findings:

1. **Vanilla MeZO diverges on HellaSwag** — eval loss climbs monotonically from R200 onward, model loses 2.5 points of accuracy by R1000. This is a new negative result: vanilla MeZO does not always converge on hard reasoning tasks, even centralized. Observed $|\hat{\rho}|$ values peak at **+159** (R360) — without clipping these outliers cumulatively drift the model.
2. **D-MeZO-N v1 rescues** — same model, same task, same hyperparameters except ρ -clip= 50 and β -decay $0.9 \rightarrow 0$ give monotonic descent (loss $2.5691 \rightarrow 2.4959$) and accuracy gain ($0.6625 \rightarrow 0.7000$, best 0.7000 reached at R800). The $\beta \rightarrow 0$ final phase produces small oscillations (R900 acc=0.6875, R1000 acc=0.7000) consistent with the β -decay corollary of Theorem 3 (§4.5): $\|v_T\|^2 \rightarrow G^2$.
3. **Federated outperforms centralized.** Federated D-MeZO-N reaches **+6.25pp accuracy** above centralized vanilla on the same Qwen3-4B / HellaSwag setup. Two compounding effects: (a) ρ -clipping + β -decay stabilization (the rescue), (b) $n = 4$ -client averaging of independent z -direction probes ($1/n$ variance-floor reduction per Theorem 2, §4.4).

This validates Theorem 3 directly: under (A4) ρ -clipping at $C = 50$, the variance bound $G^2 \leq C^2 r(H) \ell$ holds, and the iterate sequence converges linearly to the $2G^2/(3\mu)$ neighborhood. Without clipping (centralized vanilla), G^2 is unbounded and the neighborhood diverges — empirically confirmed.

5.6 Cross-lingual + cross-architecture: MathLogicQA on Qwen3.5-4B-Base

To close the universality claim, we additionally test on **MathLogicQA** (part of MERA, `ai-forever/MERA`) — 4-way symbolic-logic + arithmetic reasoning in **Russian**. This task is qualitatively different from HellaSwag: language is Russian (not English), reasoning is symbolic (not commonsense), and the suffix is a single Cyrillic letter (A/Б/В/Т) following MMLU/MERA conventions. We pair it with **Qwen3.5-4B-Base** — the hybrid linear-attention V-L architecture from §3.1 — making this the first known MeZO test on (hybrid linear-attn) $\times$ (Russian reasoning).

The data pool is MERA train (680 labelled examples, test labels are private); we split 80/20 internally to obtain 544 train / 136 val, then subsample to 500 train / 100 eval. Setup otherwise identical to §5.5.

5.6.1 v1 multi-seed falsification (fixed $C = 50$, 2-variant initial run, 3 seeds $\times$ 1000 rounds)

The single-seed Day 8 R1d hint (“D-MeZO-N v1 with fixed $C = 50$ beats vanilla by 6.0%”) was the original motivation. To avoid single-seed claims, we first ran a **2-variant** paired comparison of `vanilla` vs `dmezo_n v1` (β -decay $0.9 \rightarrow 0$ + fixed ρ -

clip $C = 50$) on the same 4-client federated setup, with 3 seeds {42, 43, 44}. Each pair shares partition and data; only the variant differs. This run **falsifies the v1 recipe**: v1 is robustly worse than vanilla.

Throughout, $\pm$ denotes the population standard deviation (ddof = 0) over $n = 3$ seeds; the sample standard deviation (ddof = 1) is $\approx 22\%$ larger and does not change any direction-of-effect statement.

Seed	vanilla loss / acc	dmezo_n loss / acc	Δ loss	Δ acc final	Δ acc peak
42	1.3747 / 0.38	1.4598 / 0.38	+0.085 (+6.2%)	0.00	0pp (both 0.39)
43	1.3432 / 0.36	1.4569 / 0.36	+0.114 (+8.5%)	0.00	+3pp (0.42 vs 0.39, dmezo_n peaks at R300)
44	1.3863 / 0.39	1.4735 / 0.39	+0.087 (+6.3%)	0.00	+2pp (0.43 vs 0.41, dmezo_n peaks at R300)
Mean $\pm$ std	1.368 $\pm$ 0.018 / 0.377 $\pm$ 0.013	1.463 $\pm$ 0.007 / 0.377 $\pm$ 0.013	+0.095 $\pm$ 0.013	0.00 $\pm$ 0.00	+1.67pp
Paired 95% CI	—	—	excludes 0	[0.000, 0.000]	within $\pm 1\sigma$

Three findings from multi-seed:

1. **Final accuracy is identical across all 3 seeds** (0.38=0.38, 0.36=0.36, 0.39=0.39). Paired Δ acc = 0.00 with bootstrap 95% CI [0.000, 0.000] — D-MeZO-N and vanilla produce **identical final classifier predictions** on the 100-example eval pool (loss differs in confidence margin only). This **falsifies** the single-seed “+1.25pp final acc” suggestion from earlier analysis.
2. **D-MeZO-N is consistently slower in loss** (3/3 seeds: vanilla wins by +6–8.5% loss, mean +7%). This is a **robust negative finding** for the MathLogicQA + IID + complete-graph regime. Mechanism (Theorem 3 in docs/theory_rigorous.md §3): the Lyapunov $V_t = (L_t - L^*) + (\eta/2)\|v_t\|^2$ contains kinetic energy; β -decay introduces transient kinetic that translates to slower loss-component convergence at fixed T=1000.
3. **Peak accuracy in trajectory is higher and earlier for D-MeZO-N on 2/3 seeds** (+2–3pp at $R \approx 300$, vs vanilla peak at $R \approx 500\text{--}700$). Mean peak boost +1.67pp, **within $\pm 1\sigma$ noise on 100-example eval (SE ± 4.9 pp)** — not significant in isolation, but the *direction is consistent across 2/3 seeds and theoretical mechanism is identified*. Practical implication: D-MeZO-N may benefit from early-stopping on accuracy metric. Pre-registration of this benefit would require fresh experiments.

5.6.2 Extended 5-variant sweep — D-MeZO-N v2 (combo) headline (3 seeds $\times$ 1000 rounds)

The v1 falsification (§5.6.1) showed the fixed $C = 50$ clip is over-tight at paper scale: on Qwen3.5-4B-Base the median $|\hat{\rho}| \approx 180$, so $C = 50$ truncates most of the useful signal. We therefore ran an **extended 5-variant** paired sweep on the same setup (Qwen3.5-4B-Base / MathLogicQA / 4 clients complete IID / 1000 rounds / $\text{lr} = 3 \cdot 10^{-7}$ / $\epsilon = 10^{-3}$ / 3 paired seeds {42,43,44}; 5 variants $\times$ 3 seeds = 15 cells, ~ 12 h Colab Blackwell). The five variants isolate each mechanism: vanilla, v1 (fixed $C = 50$), drift-only (B5 alone), adaptive_clip (B1 alone), and combo (B1+B5 = **D-MeZO-N v2**).

Variant	s42 loss/acc	s43 loss/acc	s44 loss/acc	Mean $\pm$ std loss	Mean $\pm$ std acc	Δ loss vs vanilla	Direction (3 seeds)
vanilla MeZO	1.3747 / 0.38	1.3432 / 0.36	1.3863 / 0.39	1.3681 $\pm$ 0.0182	0.3767 $\pm$ 0.0125	reference	—
D-MeZO-N v1 (fixed $C = 50$)	1.4598 / 0.38	1.4569 / 0.36	1.4735 / 0.39	1.4634 $\pm$ 0.0072	0.3767 $\pm$ 0.0125	+7.0% worse	3/3 worse (falsified)
Drift-only (B5 alone)	1.4608 / 0.38	1.4531 / 0.36	1.4537 / 0.39	1.4559 $\pm$ 0.0035	0.3767 $\pm$ 0.0125	+6.4% worse	3/3 worse (falsified)
Adaptive_clip (B1 alone)	1.2691 / 0.41	1.3135 / 0.33	1.3135 / 0.43	1.2987 $\pm$ 0.0209	0.3900 $\pm$ 0.0432	−5.1%	3/3 wins loss , acc seed-specific
D-MeZO-N v2 = combo (B1+B5, 54 resets) ★	1.2790 / 0.37	1.2951 / 0.44	1.3036 / 0.39	1.2926 $\pm$ 0.0102	0.4000 $\pm$ 0.0294	−5.5%	3/3 wins loss

Per-seed paired comparison (combo v2 vs vanilla):

Seed	vanilla loss / acc	combo v2 loss / acc	Δ loss	Δ acc
42	1.3747 / 0.38	1.2790 / 0.37	−7.0%	−1pp
43	1.3432 / 0.36	1.2951 / 0.44	−3.6%	+8pp
44	1.3863 / 0.39	1.3036 / 0.39	−6.0%	0pp
Mean	1.3681 / 0.3767	1.2926 / 0.4000	−5.5%	+2.3pp

Headline finding (loss only). D-MeZO-N v2 = combo (B1 adaptive_clip + B5 drift-reset) **robustly beats vanilla MeZO on loss**: final loss **1.293 $\pm$ 0.010 vs 1.368 $\pm$ 0.018** ($\Delta = -5.5\%$, **3/3 seeds same direction**; paired $t = 5.32$, $df=2$, $p \approx 0.034$ — **significant**). Final accuracy shows a positive but **statistically insignificant** trend (+2.3pp mean; per-seed $\{-1, +8, 0\}$ pp; paired $t \approx 0.82$, $df=2$, $p \approx 0.50$; the +8pp on s43 dominates, and on a 100-example eval pool $SE \approx 4.9$ pp) — we report it for transparency and **do not claim it as a gain**.

Four findings from the 5-variant sweep:

- v1 (fixed $C = 50$) is falsified** — 3/3 seeds worse than vanilla (+7.0% loss), with no accuracy compensation (final acc tied at 0.3767). The fixed clip is over-tight at paper scale, truncating most of the signal at the observed $|\hat{\rho}|$ scale.
- Drift-reset alone (B5) is insufficient** — 3/3 seeds worse (+6.4% loss), essentially indistinguishable from v1. Zeroing velocity does not help if the underlying clip still suppresses the signal; B5 needs a clip that lets the gradient through (B1).
- Adaptive-clip alone (B1) wins on loss but acc is seed-specific.** B1 tracks the running 95-percentile of $|\hat{\rho}|$, relaxing the clip to the data-driven range. It wins loss 3/3 (−5.1%), but its accuracy is seed-dependent (per-seed $\{0.41, 0.33, 0.43\}$, std 0.043 — the largest of any variant) and its loss std (0.021) is higher than vanilla’s (0.018): the loosened clip lets momentum overshoot, which the drift-reset is designed to catch.

4. **Combo (B1+B5 = v2) is robust on loss with the lowest std among the improving variants.** Adding drift-reset to adaptive-clip preserves the loss win (−5.5%, 3/3) **and** achieves the **lowest loss-std among the variants that improve on vanilla** (0.010 vs B1-alone 0.021 vs vanilla 0.018). Note that drift-only has a still smaller loss-std (0.0035), but this reflects it being uniformly stuck at a *worse* loss (+6.4%, 3/3 worse than vanilla), not the stability of a good solution — a tight band around a bad optimum is not a virtue. Drift-reset fired 54 times total across 3 seeds ($\approx 18/\text{seed}$): on s43 it holds the trajectory below the late-stage drift that B1-alone exhibits after R600 (B1-alone R600=1.309 $\rightarrow$ R1000=1.313; combo 1.286 $\rightarrow$ 1.295). B5 surgically zeros v_t when `eval_loss > rolling_min + 0.1`, so the Lyapunov potential component keeps contracting while kinetic energy is reset — the mechanism predicted by Theorem 3.

This makes **D-MeZO-N v2 the recommended default recipe** (adaptive clip window=50/quantile=0.95/ $\alpha = 1.3$ + drift-reset window=50/threshold=0.1 + β -decay 0.9 $\rightarrow$ 0).

5.6.3 Updated cross-task summary

Combining §5.5 (HellaSwag rescue, single-seed pending multi-seed validation), §5.6.1 (v1 falsification) and §5.6.2 (v2 headline):

Task	Vanilla MeZO	D-MeZO-N	Interpretation	Statistical status
SST-2 (Day 8 R1d)	converges	v1: 6.0% better loss	transient acceleration	single-seed (n=1)
HellaSwag	diverges (−2.5pp acc)	v1: converges (+3.75pp acc)	rescue	single-seed (pending)
MathLogicQA (3-seed)	converges (loss 1.368 $\pm$ 0.018)	v2 combo: loss 1.293 $\pm$ 0.010 (−5.5%, 3/3, $p<0.05$; acc +2.3pp n.s.)	loss win (v1 falsified, combo wins)	3 seeds, paired, $t = 5.3$

The v2 combo recipe (adaptive p-clip + drift-reset + β -decay) works as **rescue** when vanilla diverges (HellaSwag: $|\hat{\rho}|$ peaks at +159, neighborhood diverges; v1 tested, v2 rescue-regime test pending) and as a **statistically significant loss improvement** when vanilla converges (MathLogicQA, this section). Theorem 3 predicts both regimes under the same mechanism: bounded G^2 under ρ -clip $\rightarrow$ linear convergence to noise floor; without clip, unbounded $G^2 \rightarrow$ divergence.




Figure 6. Cross-domain trajectories illustrating the two D-MeZO-N regimes. (a) HellaSwag on Qwen3-4B (single-seed s42): centralized vanilla MeZO drifts upward from R200, while federated D-MeZO-N v1 descends monotonically. (b) MathLogicQA on Qwen3.5-4B-Base (single-seed view, before multi-seed validation): D-MeZO-N tracks vanilla closely. See §5.6.1 for finalized 3-seed paired CI on this task.




Figure 7. Multi-seed (n=3) validation of D-MeZO-N vs vanilla on MathLogicQA/Qwen3.5-4B-Base, federated IID complete graph. (a) Loss trajectories show mean $\pm$ std across seeds 42/43/44; vanilla consistently converges to lower final loss. (b) Accuracy trajectories show D-MeZO-N reaches higher peaks earlier in 2/3 seeds, but mean paired Δ acc final = 0.000 $\pm$ 0.000 (95% CI) — falsifying any final-accuracy advantage on this task at this scale.




Figure 8. Cross-task summary: behavior of D-MeZO-N v1 vs vanilla MeZO across three task domains. SST-2 (Day 8 R1d, single-seed): 6.0% lower final loss on worst federated cell. HellaSwag (§5.5, single-seed): +3.75pp accuracy gain in rescue regime — vanilla diverges. MathLogicQA (§5.6.1, 3-seed CI): no final-accuracy gain (paired Δ =0.0); D-MeZO-N safe-tracks vanilla in this regime, with earlier peak in 2/3 seeds.

6. Discussion

6.1 Why ring $\leq$ complete in the ZO regime?

A counter-intuitive finding: on both partition regimes the ring topology ($\rho(W) = 0.333$) consistently matches or outperforms the complete topology ($\rho(W) = 0$). In the ZO regime, very high per-step variance of $\hat{\rho}$ means that slower consensus mixing may act as an implicit regulariser.

6.2 Why is naive Nesterov incompatible with ZO?

The dual-channel noise structure of look-ahead Nesterov (probe-location and update-direction both depending on v_i) compounds the variance amplification. At $\beta = 0.9$ the look-ahead variant diverges 7 $\times$ faster than heavy-ball (R20 vs. R140).

6.3 Practical recipe and negative results on ϵ

Recipe. $\eta = 3 \cdot 10^{-7}$, $\epsilon = 10^{-3}$, ρ -clipping with $C \approx 1.3 \times \max$ observed $|\hat{\rho}|$, linear β -schedule $\beta_t = 0.9 \cdot (1 - t/T)$, doubly-stochastic mixing matrix.

Choice of constant learning rate. Classical SPSA literature (Spall 1992, §3.2) prescribes a harmonic decay schedule $\eta_t = \alpha / (t + A)^\alpha$ with $\alpha \approx 0.602$ to guarantee convergence to the true optimum θ^* . We follow Princeton MeZO (Malladi 2023) and use a **constant** η , which under Theorem 3 yields convergence only to a noise neighborhood of radius $O(G^2/\mu)$, not to θ^* . For LLM fine-tuning this is acceptable — we want a “good-enough” parameter setting, not exact convergence — and a constant η yields faster initial descent ($O(1/T)$) than harmonic decay ($O(1/T^{0.602})$). A schedule sweep over {constant, harmonic, cosine} is straightforward to add as a future ablation; the codebase exposes `MeZOConfig.lr_schedule` for this.

ϵ -schedule and ϵ -autotuner negatives. Two natural ideas for adapting the probe radius ϵ both fail empirically and are reported as honest negatives: (i) a bias-variance autotuner that selects ϵ^* by minimizing a finite-difference bias proxy chooses $\epsilon^* \in [10^{-1}, 3 \cdot 10^{-1}]$ — two orders of magnitude above the Princeton default — and loses 3–6 $\times$ in downstream loss (the proxy mismeasures the true gradient bias in bf16/fp16 arithmetic); (ii) warmup schedules $\epsilon(t)$ lose consistently across a 24-cell joint $\text{lr} \times \epsilon \times \text{variant}$ sweep (HellaSwag / Qwen3.5-4B-Base). The Princeton default $\epsilon = 10^{-3}$ remains the recommendation (tier C, §8).

6.4 Notes on acceleration, multi-direction, and batch-size variance

Note on “acceleration”. § 5.4 reports a 3 $\times$ early-stage speedup (R1b at constant $\beta = 0.9$ with clip $C = 50$) of D-MeZO-N over vanilla D-MeZO on the worst federated cell. **Our formal analysis (Theorem 3, §4.5) shows the contraction rate of D-MeZO-N, $(1 - 3\eta\mu/2)$, does not improve on the plain SGD-under-PL rate $(1 - 2\eta\mu)$ (Karimi et al. 2016).** This is consistent with Bottou–Curtis–Nocedal (2018, *SIAM Review*, Theorem 5.1): heavy-ball momentum does not yield asymptotic speedup for stochastic gradient methods with $\sigma > 0$. The 3 $\times$ speedup is therefore a **transient phenomenon** — momentum smooths variance in the early phase before the noise floor dominates. A formal **transient acceleration** bound (e.g., via multistage accelerated schemes, Aybat et al. 2019, or a Nesterov estimate sequence adapted to ZO + clip) remains an **open question**, separately tracked in `docs/theory_rigorous.md` §6.

Reviewer response: multi-direction SPSA (MD-D-MeZO-N). An external reviewer suggested replacing ρ -clipping with K -direction SPSA averaging (variance reduction at the source rather than the symptom) and using true look-ahead Nesterov to achieve the optimal $O(1/T^2)$ rate. We tested this empirically on the worst Day 5 cell (ring + Dir(0.5) Qwen3.5-4B-Base, otherwise identical to R1d) with $K = 3$. Results: final eval loss **0.1828 (K=3) vs. 0.1291 (K=1 R1d)** — **+41.6% worse on loss** — but final accuracy **0.9688 vs. 0.9563** — **+1.25pp better on acc.** K -direction averaging acts as a generalization regularizer rather than a pure speedup. The “optimal $O(1/T^2)$ Nesterov rate” claim is falsified by this result; theoretically it is also incorrect for stochastic NAG with $\sigma > 0$ (Bottou-Curtis-Nocedal 2018, Theorem 5.1). ρ -clipping and multi-direction are **complementary**, not alternatives: an ideal practitioner would use both (clip per-direction $C \cdot \sqrt{K}$). Computational cost: $2K$ forward passes per local step (vs. 2 baseline); 1.84 $\times$ wall-clock for $K = 3$ due to consensus overhead amortization.

Batch-size variance scaling (empirical: $1/\sqrt{B}$ does NOT hold). Standard CLT predicts that the per-step ρ -estimator variance shrinks as $1/B$ with mini-batch size, i.e. $\text{std} \propto 1/\sqrt{B}$. We tested this on Qwen3-0.6B / SST-2 with z fixed across 100 random batches per $B \in \{1, 2, 4, 8, 16, 32\}$ (Figure 9). The observed std plateaus at $B \geq 8$ with ratio (observed / CLT-expected) growing from $1.55 \times$ ($B=2$) to $3.43 \times$ ($B=32$). This implies the dominant noise source in MeZO is **not** data sampling — it is the choice of direction z itself (and, secondarily, low-precision arithmetic in the loss difference $L_+ - L_-$). Consequently, increasing batch size beyond ~ 8 gives no benefit, but K -direction averaging across fresh z_k DOES reduce variance (verified by `tests/test_md_mezo.py::TestKDirectionVarianceReduction`). This justifies our small-batch ($B=4-8$) recipe and motivates multi-direction extensions over larger batches.

$\hat{\rho}$ distribution vs mini-batch size B (Qwen/Qwen3-0.6B, float16, n=100/B, fixed z @seed=12345)
B=1: $\sigma=459.6245$ (expect 459.6245) B=2: $\sigma=504.9488$ (expect 325.0036) B=4: $\sigma=441.0150$ (expect 229.8123) B=8: $\sigma=261.8399$ (expect 162.5018) B=16: $\sigma=261.1344$ (expect 114.9061) B=32: $\sigma=279.0063$ (expect 81.2509)




Figure 9. Empirical distribution of the MeZO projected-gradient estimator $\hat{\rho}$ for fixed perturbation direction z across mini-batch sizes $B \in \{1, 2, 4, 8, 16, 32\}$. Setup: Qwen3-0.6B fp16 / SST-2 / $\epsilon=10^{-3}$ / 100 random batches per B. The CLT prediction $\text{std} \propto 1/\sqrt{B}$ is NOT observed: std plateaus at $B \geq 8$ (ratio observed/expected grows from $1.55\times$ at $B=2$ to $3.43\times$ at $B=32$). The dominant variance source is the choice of z , not data sampling — motivating multi-direction (vary z) over batch-scaling for variance reduction in MeZO.

6.5 Head-to-head: D-MeZO-N vs FedKSeed (Qin et al. 2024)

To address the strongest reviewer concern — the absence of direct empirical comparison with a published federated zeroth-order baseline — we implement FedKSeed (Qin et al. 2024, ICML) within our framework and run a paired head-to-head experiment on Qwen3.5-4B-Base / MathLogicQA, matching the §5.6 setup (4 clients, 3 seeds {42,43,44}, $\text{lr}=3 \cdot 10^{-7}$, $\epsilon = 10^{-3}$, IID partition) but at **500 rounds** per variant (vs 1000 in §5.6), using the D-MeZO-N v2 (combo B1+B5) recipe.

6.5.1 Algorithmic comparison

The three methods compared (see docs/fedkseed_comparison.md for full algorithmic spec):

Aspect	Vanilla D-MeZO	D-MeZO-N (ours)	FedKSeed (Qin 2024)
Per-round directions	n unique z_i	n unique z_i	1 shared z
Topology	Decentralized	Decentralized	Star (central server)
Momentum	None	Heavy-ball + β -decay	None
ρ -clipping	None	adaptive (v2: B1) + drift-reset (B5)	None
Per-edge communication	$O(d)$ (weight_avg) or $O(1)$ (update_share)	Same	$O(1)$ to server

Theoretical positioning (Theorem 2, docs/theory_rigorous.md §2). The stochastic floor $\eta C^2 r(H) \ell / (\mu n)$ achieves $1/n$ federated speedup only with **independent direction sampling per client**. FedKSeed uses one shared z per round, so its noise floor reduces only the data-sampling component by $1/n$ — direction noise σ_z^2 is not reduced. D-MeZO and D-MeZO-N reduce both, predicting lower variance floors.

6.5.2 Empirical results

The experiment was run with scripts/head_to_head_fedkseed.py on the same Colab Blackwell setup as §5.6. **The table below reports finalized 3-seed mean $\pm$ std** (paired analysis: same seed pairs vanilla, dmezo_n v2, fedkseed cells; data

partition is deterministic per seed). All numbers are at 500 rounds/variant.

Variant	Final loss (mean $\pm$ std)	Final accuracy (mean $\pm$ std)	Per-seed loss	Comm bytes/round
Vanilla D-MeZO	1.4630 $\pm$ 0.0229	0.3733 $\pm$ 0.0309	1.4945 / 1.4405 / 1.4542	64 B (update_share)
D-MeZO-N v2 (ours)	1.3341 $\pm$ 0.0142 ★	0.3900 $\pm$ 0.0374	1.3540 / 1.3228 / 1.3254	64 B (update_share)
FedKSeed (Qin 2024)	1.4656 $\pm$ 0.0225	0.3367 $\pm$ 0.0330	1.4668 / 1.4925 / 1.4375	~64 B (4 scalars to server)

Loss ranking (the headline metric). D-MeZO-N v2 beats **both** baselines on loss **3/3 seeds**: it beats FedKSeed on loss 3/3 seeds and beats vanilla on loss 3/3 seeds. FedKSeed and vanilla are statistically indistinguishable on loss (1.4656 vs 1.4630).

Paired analyses on accuracy (bootstrap 95 % CI, $n=3$):

- Δ_{acc} (D-MeZO-N v2 vs vanilla): **+0.0167**, CI [+0.0000, +0.0400] — **not significant**.
- Δ_{acc} (D-MeZO-N v2 vs FedKSeed): **+0.0533**, CI [-0.0200, +0.1400] — **not significant**.
- Δ_{acc} (FedKSeed vs vanilla): **-0.0367**, CI [-0.1000, +0.0300] — **not significant**.

All accuracy comparisons are not significant at $n = 3$ (the 100-example eval pool gives $\text{SE} \approx 4.9\text{pp}$); the robust signal is on loss. (With $n = 3$ seeds the paired bootstrap has only 10 distinct resamples; we report these intervals as conservative ranges rather than calibrated CIs.)




Figure 10. Head-to-head comparison of D-MeZO-N v2, vanilla MeZO and FedKSeed on Qwen3.5-4B-Base / MathLogicQA (3 seeds, 500 rounds, 4 clients IID): side-by-side loss and accuracy trajectories with shaded $\pm 1\sigma$ bands across seeds. D-MeZO-N v2 separates from both baselines on loss in all three seeds.

6.5.3 Interpretation

D-MeZO-N v2 is the only one of the three methods that improves on the vanilla loss, and it does so on every seed. FedKSeed, by contrast, tracks vanilla on loss (no momentum/clip to stabilize the trajectory) and trends slightly below vanilla on accuracy (-3.7pp mean, not significant). This is directionally consistent with the Theorem 2 positioning (§6.5.1): D-MeZO-N’s independent per-client z_i reduces **both** the data- and direction-noise components by $1/n$, whereas FedKSeed’s single shared z per round reduces only the data-sampling component — predicting a lower variance floor (and hence lower loss) for D-MeZO-N. The accuracy comparisons are inconclusive at this scale and we do not claim an accuracy advantage.

Important caveat (parity tuning). FedKSeed was run at its default $K = 4096$ hyperparameters **without a dedicated $\text{lr} \times \beta$ grid search**; a parity-tuned comparison (matching the per-method hyperparameter search budget) is future work. The

comparison is also at **500 rounds** (not the 1000 rounds of §5.6), so the absolute losses are higher than the §5.6 headline. These two caveats mean the loss gap should be read as “v2 with our recipe beats FedKSeed-at-default”, not as a tuned head-to-head.

Reproducibility caveat: Our FedKSeed implementation is faithful to Qin et al. § 3 algorithm description (shared per-round seed, scalar aggregation, no momentum/clip) but runs in our codebase, not their official repository (github.com/alibaba/FederatedScope/tree/FedKSeed). For one cell we plan a cross-check using their original code as an integrity test; methodology described in `docs/fedkseed_comparison.md`.

6.6 Local ablation of proposed improvements (B5 drift-reset, B1 adaptive-clip, D2 DP-MeZO)

Beyond the main D-MeZO-N v1 recipe (β -decay $0.9 \rightarrow 0$ + fixed ρ -clip $C = 50$), we explore three further improvements derived from the analysis in §5 and §6:

- **B5 (drift-reset).** Zero out the velocity buffer if eval loss has risen by more than τ over a sliding window of W rounds. Motivated by R1b late-stage drift (§5.4): momentum buffer accumulates noise after convergence and pushes loss back up. Reset is a cheap and explicit fix that bypasses the need for β -decay.
- **B1 (adaptive ρ -clip).** Replace fixed $C = 50$ with $C_t = \alpha \cdot \text{quantile}_{0.95}(\{|\hat{\rho}_{t-W}|, \dots, |\hat{\rho}_{t-1}|\})$. The fixed $C = 50$ was chosen on §5.4 SST-2 ablations; running quantile adapts to the per-task ρ distribution (which differs by 2-10 $\times$ across tasks per §6.7 diagnostics).
- **D2 (DP-MeZO).** Add $\mathcal{N}(0, \sigma_{\text{DP}}^2)$ noise to the (clipped) projected gradient. The clip threshold C provides L2-sensitivity bound $\Delta = C$, giving Gaussian-mechanism (ϵ, δ) -DP: $\epsilon = C\sqrt{2\ln(1.25/\delta)}/\sigma_{\text{DP}}$ (Dwork-Roth 2014).

These were implemented as opt-in flags (default-disabled, full backwards compat — 18 unit tests in `tests/test_improvements.py`). We tested locally on **Qwen3.5-0.8B / SST-2 / 200 rounds / 4 clients IID / 2 seeds** for fast iteration (~36 min wall-clock on RTX 5070 Ti Blackwell with newly-installed `triton-windows` + `flash-linear-attention`).

6.6.1 Results

Variant	Final loss (mean $\pm$ std)	Final acc	Best acc	Drift resets	Δ loss vs vanilla
Vanilla D-MeZO	0.241 $\pm$ 0.011	0.925	0.950	0	—
D-MeZO-N v1 (fixed C, β -decay)	0.807 $\pm$ 0.017	0.925	0.9375	0	+0.566
+ B5 drift-reset	0.805 $\pm$ 0.002	0.9375	0.9375	6 (3/seed)	+0.564
+ B1 adaptive-clip ($\alpha = 1.3$)	0.510 $\pm$ 0.120	0.750 $\pm$ 0.050	0.925	0	+0.270
+ D2 DP ($\sigma = 0.5$, $\epsilon = 378$)	0.803 $\pm$ 0.026	0.925	0.950	0	+0.562

Paired bootstrap 95 % CI on Δ_{acc} (improvement vs vanilla) (with $n = 2$ seeds the paired bootstrap has only 3 distinct resamples; we report these intervals as conservative ranges rather than calibrated CIs):

Improvement	Δ_{loss}	Δ_{acc}	95 % CI on Δ_{acc}	Significant?
B5 drift-reset	+0.565	+0.013	[-0.025, +0.050]	NO
B1 adaptive-clip	+0.270	-0.175	[-0.200, -0.150]	YES (negative)
D2 DP ($\sigma = 0.5$)	+0.562	0.000	[0.000, 0.000]	NO

6.6.2 Three findings

1. Vanilla wins on this convergent task (consistent with §5.6.1). SST-2 with Qwen3.5-0.8B converges very quickly: vanilla reaches loss 0.2405 in 200 rounds, all D-MeZO-N variants reach ~ 0.8073 ($\approx 3.4\times$ worse; precisely $0.8073/0.2405 = 3.36\times$). This **replicates the §5.6.1 multi-seed MathLogicQA pattern**: on tasks where vanilla converges quickly, D-MeZO-N’s momentum + clip recipe is over-engineered. **D-MeZO-N’s value lies in the rescue regime** (§5.5 HellaSwag, where vanilla diverges), not in routine convergent training.

2. Adaptive-clip paradox. The adaptive threshold settles at $\sim 400\text{--}500$ during training — **8-10 $\times$ higher than fixed $C = 50$** . This confirms our hypothesis that the fixed clip used in D-MeZO-N v1 is over-aggressive (cutting too much signal). With the looser clip, loss converges faster (0.51 vs 0.81 for fixed- C v1, -37% improvement), but **accuracy degrades by 17.5 pp** with 95 % CI [-0.200, -0.150] **excluding zero**. The trajectory shows quick initial descent followed by oscillation — classic momentum overshoot once the clip stops attenuating large velocity updates. **Inter-seed variance** for adaptive-clip is also high (loss std 0.12 vs 0.011 for vanilla — $\sim 10\times$ more sensitive to the random seed). Practical implication: adaptive-clip cannot be used alone; it needs combining with B5 drift-reset (a velocity-zeroing safety net) or with a smaller $\alpha < 1.0$.

3. Drift-reset and DP work mechanically but show no measurable benefit at 200 rounds. Drift-reset fired $3\times/\text{seed}$ (window=50, threshold=0.1) but the trajectory is nearly identical to D-MeZO-N v1 — 200 rounds is too short for the late-stage drift this mechanism is designed to fix (R1b drift was visible only after R300 in §5.4). DP at $\sigma = 0.5$ has no measurable degradation, demonstrating that the mechanism is composable, but only provides $\varepsilon \approx 378$ — far from the $\varepsilon \leq 10$ range needed for meaningful privacy. A proper σ -sweep at $\sigma \in \{5, 10, 19\}$ to map the privacy/utility frontier is left for future work.




Figure 11. Local ablation of D-MeZO-N improvements (B5 drift-reset, B1 adaptive-clip, D2 DP-MeZO) on Qwen3.5-0.8B / SST-2 / 200 rounds / 4 clients IID / 2 seeds. (a) Loss trajectories: vanilla converges fastest; adaptive-clip is the only D-MeZO-N variant approaching vanilla’s loss, but at the cost of (b) accuracy: adaptive-clip drops to 0.75 final accuracy (down from 0.93 init), with CI excluding zero — significant negative effect. Drift-reset and DP are essentially indistinguishable from D-MeZO-N v1 at this horizon. The shaded bands are ± 1 standard deviation across the 2 seeds.

6.6.3 Implications

The local ablation is **another honest negative result for D-MeZO-N improvements on convergent tasks**, mirroring §5.6.1. Three actionable conclusions:

1. **B1 adaptive-clip exposes the true bottleneck:** D-MeZO-N's fixed $C = 50$ is suboptimal across tasks (the empirically-discovered 0.95-quantile is $\sim 8\times$ higher), but a naive loosening trades loss-speed for trajectory-stability. The fix is **combining B1 with B5** (or with $\alpha < 1.0$) — natural next experiment, not yet run.
2. **B5 drift-reset is currently a no-op at 200 rounds.** Its value is preventing the R1b-style late-drift catastrophe, which requires $T \geq 500$ to observe. We re-classify B5 as **stabilization insurance**, not a performance improvement.
3. **D2 DP-MeZO has correct mechanics but needs proper privacy budget sweep.** At $\sigma = 0.5$ we get $\varepsilon = 378$ (no privacy) with $\sim$ zero utility cost — the **mechanism composes cleanly**. The right next experiment is $\sigma \in \{5, 10, 19, 50\}$ to find the privacy/utility frontier. This is a paper-changing direction: **first privacy-preserving decentralized federated ZO optimizer** if we hit $\varepsilon \leq 10$ with acceptable utility loss.

The implementation, smoke tests, and analysis script are reproducible via `scripts/local_test_improvements.py` and `scripts/analyze_local_test.py`. See `docs/upgrade_roadmap.md` for the prioritized next steps.

6.6.4 Following the adaptive-clip lead: B1+B5 combo recovers vanilla parity

Section 6.6.2 identified that adaptive-clip alone trades better loss for worse accuracy due to momentum overshoot. The natural fix is to **combine B1 with B5** — adaptive clip relaxes the over-tight $C = 50$ and lets the gradient signal through, while drift-reset zeros the velocity when overshoot causes loss to rise. We test this on a harder task than SST-2: **Qwen3.5-0.8B / MathLogicQA (Russian 4-way symbolic logic) / 200 rounds / 4 clients IID / 2 seeds**.

Variant	Final loss	Final acc	Best acc	Resets	Δ loss vs vanilla	Paired 95% CI on Δ acc
Vanilla D-MeZO	1.4738 ± 0.005	0.3750	0.4125	0	—	(reference)
D-MeZO-N v1 (fixed C)	2.0808 ± 0.017	0.2625	0.3250	0	+0.607	[-0.137, -0.087] (worse)
+ B5 drift-reset only	2.0569 ± 0.028	0.2625	0.3250	6	+0.583	[-0.125, -0.100] (worse)
+ B1 adaptive-clip only	1.4810 ± 0.006	0.3250	0.3875	0	+0.007	[-0.075, -0.025] (worse, small)
+ B1+B5 combo (new)	1.4735 ± 0.026	0.3250 ± 0.025	0.3875	6	-0.0004 (tie)	[-0.100, +0.000] (tie, includes 0)

Three statistical findings:

1. **Combo achieves loss-parity with vanilla** ($\Delta_{\text{loss}} = -0.0004$, essentially zero). This is a **29% improvement** over D-MeZO-N v1's loss of 2.08.
2. **Combo's accuracy difference from vanilla is statistically indistinguishable from zero** (paired 95 % CI includes 0). This is a **6.25 pp improvement** over D-MeZO-N v1's accuracy of 0.2625, statistically significant against v1 (CI [+0.025, +0.100] excludes zero).

3. **The drift-reset triggers 6 times across 2 seeds** (3 per seed) — confirming it actively catches the adaptive-clip overshoot. On seed=43 specifically, combo reaches **loss 1.45 / acc 0.35**, matching vanilla’s (loss 1.48 / acc 0.35) on the same seed exactly.

Mechanism (Theorem 3 framing). D-MeZO-N v1’s fixed $C = 50$ over-attenuates the gradient signal (adaptive-clip discovers the effective threshold is 8-10 $\times$ higher). Without B5, this lets the unleashed momentum accumulate kinetic energy that pushes loss up after the initial descent — the SST-2 paradox of §6.6.2. With B5, the velocity buffer is zeroed when this overshoot is detected (loss rise > 0.1 over a 50-round window), letting the Lyapunov $V_t = (L - L^*) + (\eta/2)\|v\|^2$ contract on its loss component without kinetic-energy buildup. The empirical result confirms the theoretical mechanism.




Figure 12. B1+B5 combo recovers vanilla parity on MathLogicQA. (a) Loss trajectories: vanilla (red), B1 adaptive-clip (purple), and B1+B5 combo (cyan) all converge to ~ 1.47 final loss. D-MeZO-N v1 (blue) and B5 drift-reset alone (green) are stuck at ~ 2.05 loss. (b) Accuracy trajectories: vanilla maintains highest acc; combo and adaptive-clip reach ~ 0.32 - 0.35 (close to vanilla 0.37); drift alone and v1 remain near random (0.25-0.275). Setup: Qwen3.5-0.8B / 4 clients IID / 200 rounds / 2 seeds.

Re-framing the D-MeZO-N recipe. Combining the observations from §6.6.2 (SST-2 vanilla wins, adaptive-clip paradox) and §6.6.4 (MathLogicQA combo achieves parity):

- **D-MeZO-N v1** (fixed $C = 50 + \beta\text{-decay } 0.9 \rightarrow 0$) — empirically over-tight clip; 3.4 $\times$ worse loss than vanilla on convergent tasks; *as documented in §5.6.1 multi-seed*.
- **D-MeZO-N v2 (proposed, with B1+B5 combo)** — adaptive clip threshold $\alpha = 1.3 \times \text{quantile}_{0.95}(\{|\hat{\rho}|\}_{\text{recent}})$ + drift-reset (window=50, threshold=0.1) + β -decay. **Recovers vanilla parity on MathLogicQA; statistically significant improvement over v1.**

This is the first concrete improvement to the D-MeZO-N recipe since the original v1 formulation in §5.4. We recommend **v2 as the new default recipe** for D-MeZO-N going forward, pending replication at paper scale (Qwen3.5-4B-Base, see [docs/upgrade_roadmap.md §A.2](#)).

Caveats and remaining work: - Only $n = 2$ seeds — paper-scale validation needs $n = 3$ multi-seed at Qwen3.5-4B (matching §5.6.1). - Combo introduces 5 new hyperparameters (α , ac_window , ac_quantile , drift_window , drift_threshold). Robustness under perturbation not tested. - Tested only on local Qwen3.5-0.8B (much smaller than paper-scale Qwen3.5-4B). The benefit may differ at scale. - D-MeZO-N v2 **must** also be tested on **rescue regime** (HellaSwag §5.5) to confirm it doesn’t break the rescue ability, since unclipped momentum could re-introduce divergence on tasks where vanilla diverges.

Scale-up validation (completed at paper scale). Both follow-up experiments planned here were subsequently run or scheduled. The paper-scale replication (Qwen3.5-4B-Base / MathLogicQA / 3 seeds $\times$ 1000 rounds, ~ 12 h Colab Blackwell) is reported as the headline result in §5.6.2: D-MeZO-N v2 (combo) beats vanilla on loss 3/3 seeds ($\Delta = -5.5\%$, paired $t =$

5.3, $p \approx 0.034$) — the small-scale parity finding of this section strengthens into a win at scale, making **v2 the recommended default recipe**. The rescue-regime stress test (Qwen3-4B / HellaSwag / 3 seeds) remains pending (§7; tier D1 in §8): B1 relaxes the over-tight $C = 50$ to ~ 400 , which could re-introduce divergence on tasks where vanilla diverges; B5 drift-reset is designed to catch this, but the safety margin at large $|\hat{\rho}|$ peaks (+159 observed in §5.5) is empirically untested.

6.7 Privacy-preserving D-MeZO-N (DP-MeZO σ -sweep)

We extend D-MeZO-N v1 with Gaussian noise on the clipped projected gradient, transforming each MeZO step into one application of the **Gaussian mechanism** (Dwork-Roth 2014). The clip threshold C naturally serves as the L2 sensitivity bound, so adding $\xi_t \sim \mathcal{N}(0, \sigma^2)$ to $\tilde{\rho}_t = \text{clip}(\hat{\rho}_t, \pm C)$ gives (ε_1, δ) -DP per round with $\varepsilon_1 = C \sqrt{2 \ln(1.25/\delta)}/\sigma$.

Theoretical analysis is in docs/theory_rigorous.md §6.5 (Theorem 4). Key result:

$$\mathbb{E}[V_T] \leq (1 - \frac{3\eta\mu}{2})^T V_0 + \frac{2(C^2 r(H) + \sigma^2 d)\ell}{3\mu}$$

where the $\sigma^2 \cdot d$ term reflects an important theoretical fact: **DP-noise breaks the Malladi $r(H)$ -substitution trick** because ξz is isotropic in parameter space (no alignment with ∇L), so the noise contribution scales with **full d** rather than $r(H) \ll d$, while the clip-noise contributes $C^2 r(H)\ell$ (so the bound reduces exactly to Theorem 3 at $\sigma = 0$). The theoretical crossover where DP-noise dominates ZO-noise is $\sigma_{\text{crossover}} \sim C \sqrt{r(H)/d}$ which is very small (≈ 0.02 for our setup); but empirically the floor bound is loose at $T = 200$ rounds.

6.7.1 σ -sweep setup

Experimental setup (matches docs/dp_sigma_sweep_plan.md): - Model: Qwen/Qwen3.5-0.8B (hybrid linear-attn; smaller model for faster sweep — σ -sensitivity should be approximately model-size-invariant for the per-round bound) - Task: MathLogicQA (4-way Russian symbolic logic) - Federation: 4 clients, IID, complete topology, weight_avg consensus - 2 seeds $\times$ 200 rounds $\times$ 8 variants (6 σ values + vanilla + dmezo_n no-DP baseline) = 16 cells

Privacy budget table ($C = 50$, $\delta = 10^{-3}$):

σ	ε per round	Privacy class	Predicted utility
0.5	378	No privacy (reference)	$\approx$ no-DP D-MeZO-N
2.0	94	Trivial	$\approx$ no-DP D-MeZO-N
5.0	38	Trivial	Slight degradation
10.0	19	Weak	Moderate degradation
19.0	10	Medium ★ paper threshold	Key data point
50.0	4	Medium-strong	Significant degradation expected

6.7.2 Empirical results (Colab Blackwell, 119 min wall-clock)

The DP σ -sweep was executed on Colab Pro+ Blackwell on 2026-05-20 (16 cells: 8 variants $\times$ 2 seeds, ~ 7.5 min per cell).

Variant	σ	ε per round	Final loss (mean $\pm$ std)	Final acc (mean $\pm$ std)	$\Delta_{\text{loss vs no-DP D-MeZO-N}}$
Vanilla D-MeZO	—	∞ (no DP)	1.4698 $\pm$ 0.001	0.310 $\pm$ 0.030	— (lower bound, no momentum/clip)
D-MeZO-N v1 (no DP)	—	∞ (no DP)	1.7854 $\pm$ 0.018	0.265 $\pm$ 0.025	0% (reference)
+ DP, $\sigma = 0.5$	0.5	378	1.9075 $\pm$ 0.054	0.255 $\pm$ 0.025	+6.8%
+ DP, $\sigma = 2.0$	2.0	94	1.8933 $\pm$ 0.065	0.255 $\pm$ 0.015	+6.0%
+ DP, $\sigma = 5.0$	5.0	38	1.8828 $\pm$ 0.098	0.250 $\pm$ 0.030	+5.5%
+ DP, $\sigma = 10.0$	10.0	19	1.9068 $\pm$ 0.076	0.230 $\pm$ 0.010	+6.8%
+ DP, $\sigma = 19.0$	19.0	★ 10	1.8967 $\pm$ 0.093	0.265 $\pm$ 0.035	+6.2%
+ DP, $\sigma = 50.0$	50.0	4	1.9116 $\pm$ 0.036	0.275 $\pm$ 0.045	+7.1%

The full privacy/utility frontier (final eval loss and accuracy as functions of the per-round privacy budget ε on a log scale, with no-DP baselines and ± 1 std error bars across seeds) is visualized by `scripts/sweep_dp_sigma.py`; the table above contains the complete underlying data. The frontier is **statistically flat across all σ values** — utility loss vs no-DP D-MeZO-N is essentially constant at 5.5–7.1%, regardless of ε .

Data availability. The raw σ -sweep JSONs reside in the Colab run archive (`dmezo_runs/` on Drive; this sweep was executed Colab-only on 2026-05-20) and are not committed to the repository; the sweep is fully regenerable via `scripts/sweep_dp_sigma.py`. The repository’s `experiments/diagnostics/` covers a single DP configuration ($\sigma = 0.5$).

6.7.3 Main finding: DP is essentially free for D-MeZO-N at this scale

The empirical privacy/utility frontier is statistically flat across $\sigma \in [0.5, 50]$ — i.e., the entire tested ε range from 378 (no privacy) down to 4 (medium-strong privacy) gives **the same loss and the same accuracy as no-DP D-MeZO-N v1, within seed noise**. All confidence intervals overlap; the loss values cluster around 1.90 ± 0.07 and the accuracy around 0.25 ± 0.03 .

Per-round $\varepsilon = 10$ ($\sigma = 19$) finding: DP-MeZO-N achieves loss 1.8967 vs no-DP 1.7854 (+6.2%) and accuracy 0.265 vs 0.265 (identical, within noise). This establishes:

The first decentralized federated zeroth-order optimizer with formal (ε, δ) -DP guarantee at $\varepsilon = 10$ on LLM fine-tuning, with only $\sim 6\%$ utility cost vs the same-framework no-DP baseline.

The mechanism is elegant: ρ -clipping (the velocity-stabilization component of D-MeZO-N v1, originally introduced to control momentum overshoot in §5.4) provides the **natural L2 sensitivity bound $\Delta = C$ for the Gaussian mechanism**. No additional clipping or per-example gradient norm computation is required, in contrast to DP-SGD which needs explicit per-sample gradient clipping (Abadi et al. 2016). The single scalar $\hat{\rho}$ already carries the Gaussian mechanism’s full DP guarantee.

6.7.4 Why is the theoretical noise floor not observed?

Theorem 4 (§6.5 of `docs/theory_rigorous.md`) predicts steady-state noise floor $\frac{2(C^2 r(H) + \sigma^2 d)\ell}{3\mu}$, with the $\sigma^2 d$ term dominating for $\sigma > C\sqrt{r(H)/d} \approx 0.02$. With $d \sim 0.85 \cdot 10^9$ for Qwen3.5-0.8B, the predicted noise floor at $\sigma = 19$ should be catastrophic (theoretically $\sim 5 \times 10^4 \times$ the no-DP floor).

The empirical observation that **utility is essentially constant** instead means one of three things:

1. **Finite-horizon effect** — at $T = 200$ rounds, the optimization has not converged to the steady-state floor; the $(1 - 3\eta\mu/2)^T V_0$ transient term still dominates over the noise floor. This is consistent with Theorem 4a: for small enough T , the bound is dominated by transient decay, not noise.
2. **Effective d is much smaller than total parameter count**. Only trainable parameters contribute (vision tower frozen; only text decoder participates), and within those, the alignment between z and ∇L projects most of the random direction onto irrelevant subspaces. Effective d in our setup may be $\sim 10^6$ rather than 10^9 .
3. **Discretization** — at finite step size $\eta = 3 \times 10^{-7}$ and small T , the SDE-style steady-state analysis is not yet a tight predictor.

This loose bound is a **feature, not a bug**: real-world deployments with $T \leq 1000$ rounds will likely also enjoy this gap, meaning DP can be applied liberally at $\varepsilon \leq 10$ without measurable utility cost.

6.7.5 Alternative narratives (not supported by data)

For honesty: we pre-registered three outcome narratives (A/B/C) in `docs/dp_sigma_sweep_plan.md` before running the sweep. The data unambiguously supports Scenario A (above). The discarded narratives:

- **(B) “Smooth degradation with σ ”** — Empirically rejected: the frontier is statistically flat, not monotonically decreasing.
- **(C) “Catastrophic collapse at $\sigma \geq 19$ requiring K-direction averaging”** — Empirically rejected: no collapse observed up to $\sigma = 50$.

6.7.6 Composition caveat (honest)

Per-round ε reported above is for **one round** of training. **T-round composition** is fundamentally harder:

- **Basic composition** (Dwork-Roth Theorem 3.16): $\varepsilon_T = T\varepsilon_1$. For $T = 200, \varepsilon_1 = 10$: $\varepsilon_T = 2000$ (useless).
- **Advanced composition** (Dwork-Rothblum-Vadhan 2010): $\varepsilon_T = O(\sqrt{T \ln(1/\delta')}) \cdot \varepsilon_1 + O(T\varepsilon_1(e^{\varepsilon_1} - 1))$. The second term is catastrophic for $\varepsilon_1 > 1$; advanced/RDP composition yields meaningful savings **only in the $\varepsilon_1 \ll 1$ regime**, where $e^{\varepsilon_1} - 1 \approx \varepsilon_1$.
- **RDP / moments accountant** (Mironov 2017, Abadi 2016): the Gaussian mechanism with L2-sensitivity C and noise σ is $(\alpha, \frac{\alpha C^2}{2\sigma^2})$ -RDP per round (the sensitivity factor C^2 must be included — the bare $\alpha/(2\sigma^2)$ coefficient is for unit-sensitivity), giving $(\alpha, \frac{T\alpha C^2}{2\sigma^2})$ -RDP after T rounds, still $O(\sqrt{T})$ in the converted (ε, δ) budget for Gaussian.
- **Subsampling amplification** (Abadi 2016): reduces per-step ε by batch fraction q , but our setup uses full-batch MeZO per round (no subsampling).

Paper position: Report per-round ε (standard convention for one-shot federated fine-tuning); explicitly note T-round composition as a limitation. Achieving a useful training-level budget requires per-round $\varepsilon_1 < 1$ (i.e. larger σ), which is the regime a full RDP-accountant treatment must target (future work). Adding mini-batch subsampling per round is a straightforward extension that would amplify privacy.

7. Limitations and Future Work

Empirical.

- **Multi-seed coverage.** Day 5 SST-2 grid has $n = 2$ seeds; HellaSwag (§5.5) is single-seed (multi-seed rescue validation is configured but pending Colab budget — `validate_dmezo_n_rescue_multiseed.py`, Section 19 in Colab notebook). MathLogicQA (§5.6) is **finalized**: the full 5-variant $\times$ 3-seed sweep (15 cells, §5.6.2) is complete, and the headline D-MeZO-N v2 loss improvement over vanilla ($\Delta\text{loss} = -5.5\%$, 3/3 same direction, paired $t = 5.3$, $p < 0.05$) is multi-seed validated. The earlier single-seed “+3.75pp peak-acc” hint did **not** survive multi-seed checking: final accuracy is identical across seeds for vanilla/v1/drift-only (§5.6.1), and the v2 accuracy trend (+2.3pp mean) is statistically insignificant ($p \approx 0.5$, $\text{SE} \pm 4.9\text{pp}$ on the 100-example pool) and reported for transparency only. Conservative reading: on convergent tasks D-MeZO-N v2 delivers a robust **loss** improvement and **safe-tracks** vanilla on accuracy (no divergence); accuracy gains beyond seed noise are not claimed at this scale.
- **Head-to-head vs FedKSeed is at default hyperparameters only.** §6.5 reports a 3-seed paired comparison in which D-MeZO-N v2 beats FedKSeed on loss 3/3 seeds, but FedKSeed was run at its default $K = 4096$ without a dedicated $\text{lr} \times \beta$ grid search, and only at 500 rounds. A parity-tuned comparison (and comparisons vs Ferret / FedZeN, whose integration is non-trivial due to different topology and LoRA-vs-full-parameter assumptions) remains future work.
- **Scale.** Tested at $n = 4$ clients and 4B parameters. Real federated deployments scale to 100+ clients and 8B+ models. One scale-up experiment (e.g., Qwen3-8B or $n = 8$ clients) would address this; planned in `docs/upgrade_roadmap.md` §B.3.
- **Task coverage.** All experiments are multi-choice classification (SST-2 binary; BoolQ binary; HellaSwag 4-way; MathLogicQA 4-way). No generative tasks (SAMSum, GSM8K, free-form QA) tested — the generative loss landscape may differ qualitatively in $r(H)$ and Hessian conditioning.
- **Post-hoc tuning.** D-MeZO-N v1 hyperparameters (β -decay $0.9 \rightarrow 0$, clip $C = 50$) were chosen via ablation on SST-2 (Day 8 R1d). HellaSwag (§5.5) and MathLogicQA (§5.6) used the **same** hyperparameters **without re-tuning** — we argue this as a universality test, but ideally would have been pre-registered before running.

Theoretical.

- **Theorem 3** establishes **stability and convergence to a noise floor** under PL + heavy-ball + β -decay, at contraction rate $(1 - 3\eta\mu/2)$ which **does not improve on the plain SGD-under-PL rate** $(1 - 2\eta\mu)$ (Karimi et al. 2016). **No asymptotic momentum acceleration is claimed** — the contribution is that momentum can be made *safe* (non-divergent) in the high-variance ZO regime. This is consistent with Bottou–Curtis–Nocedal 2018 (Thm 5.1).
- **Empirical $3\times$ transient speedup** (Day 8 R1b at constant $\beta = 0.9$ before R300) is **not yet formally explained**. A transient acceleration analysis (estimate sequence adapted to ZO + ρ -clip) is an open question (`docs/theory_rigorous.md` §6).
- **Full decentralized extension of Theorem 3** (combining $\rho_W < 1$ mixing with heavy-ball under PL) is incomplete; the centralized version (this paper) plus convex decentralized (Theorem 1) cover two of the three needed corners — the decentralized PL + heavy-ball case is an open question (`docs/theory_rigorous.md` §6).
- **Look-ahead Nesterov variant** has dual-channel noise structure (probe + update both depend on v_t) — empirically diverges $7\times$ faster than heavy-ball (R20 vs R140). No theoretical bound derived; matches qualitative reasoning but not formally closed.
- **PL constant μ for LLMs is unproven globally** — locally on overparameterized trajectory it is plausible (Liu-Zhu-Belkin 2022) but we use it as assumption.

8. Calibrated summary of achievements

Tiers ranked by evidence strength, separating what has been **solidly demonstrated**, what is **promising but preliminary**, what has been **falsified**, and what is **pending**. Each claim cites the relevant experimental section. (Within tier A, rows are ordered by evidence strength, not by label number.)

A. Solidly demonstrated (multi-seed or task-replicated, paper-ready)

Claim	Evidence	Section
A1. Federated MeZO works on hybrid linear-attention LLMs (Qwen3.5-4B-Base)	Day 1 + 2×2 cross-arch grid	§5.1, §5.2
A2. 2D-torus, ring, complete topologies all converge; partition-tax of Dir(0.5) non-IID is <13%	Day 5 2×2 grid	§5.2
A3. Two new convergence theorems with full Lyapunov proofs: T3 (PL+heavy-ball+ β -decay+clip; first stability guarantee for ZO heavy-ball in the PL regime) and T4 (DP extension)	Closed proofs in <code>docs/theory_rigorous.md</code>	§4.5, §4.6
A7. 🌟 D-MeZO-N v2 = combo (B1+B5) beats vanilla MeZO on loss at paper-scale (Qwen3.5-4B-Base / MathLogicQA / 3 seeds paired): $\Delta\text{loss} = -5.5\%$ (3/3 same direction, paired $t = 5.3$, $p < 0.05$); lowest loss-std (0.010) among the variants that improve on vanilla (drift-only's 0.0035 is tighter but stuck at a worse loss). Acc +2.3pp reported for transparency, n.s.	5-variant × 3-seed multi-seed run	§5.6.2
A8. D-MeZO-N v2 beats FedKSeed (Qin 2024) on loss 3/3 seeds in a paired head-to-head (FedKSeed at default $K = 4096$, 500 rounds)	3-method × 3-seed run	§6.5
A4. Per-round ($\epsilon=10$, $\delta=10^{-3}$)-DP at ~6% utility cost on Qwen3.5-0.8B / MathLogicQA	8 σ values × 2 seeds = 16 cells, frontier flat across $\sigma \in [0.5, 50]$	§6.7
A5. Communication cost: 1 scalar + 1 seed per client per round (~16 bytes), vs FedAvg's $O(d)$	Algorithmic	§4
A6. ρ -clip threshold C provides natural L2 sensitivity bound for Gaussian DP mechanism (no extra mechanism needed)	Insight + empirical confirmation	§6.7

B. Promising but preliminary (1-2 seeds, replication pending)

Claim	Evidence	Caveat
B1. D-MeZO-N v1 rescues HellaSwag where vanilla diverges (+3.75pp acc on Qwen3-4B)	n=1 seed	Multi-seed pending; should not be over-claimed
B2. D-MeZO-N v2 (B1+B5 combo) achieves vanilla parity on Qwen3.5-0.8B / MathLogicQA (small-scale ablation)	n=2 seeds; paired CI [-0.10, 0.00] on Δacc includes zero	Superseded by the paper-scale A7 result (§5.6.2)

C. Empirically falsified (honest negatives — kept in paper)

Original claim	Falsification	Section
C1. “D-MeZO-N v1 (fixed $C = 50$) beats vanilla on MathLogicQA”	3-seed paired: 3/3 worse (+7.0% loss)	§5.6.1, §5.6.2 (multi-seed)
C1b. “Drift-reset (B5) alone is sufficient”	3-seed paired: 3/3 worse (+6.4% loss)	§5.6.2
C2. “True-Nesterov look-ahead accelerates”	Diverges 7× faster than heavy-ball (R20 vs R140)	§5.4 (Day 6b)
C3. “K=3 multi-direction strictly improves”	Trade-off: loss +41.6% worse, acc +1.25pp better	§6.4 (MD ablation)
C4. “Adaptive $\epsilon(t)$ autotuner beats Princeton $\epsilon=10^{-3}$ ”	Loses by 3-6× in downstream	§6.3 (ϵ autotuner)
C5. “ $O(1/T^2)$ Nesterov rate achievable”	Bottou-Curtis-Nocedal 2018 Theorem 5.1: momentum does not accelerate asymptotically for stochastic non-convex with $\sigma>0$	§6.4 theoretical note

D. Pending validation (in flight or planned)

Task	Status	ETA
D1. Combo on HellaSwag rescue regime / Qwen3-4B / 3 seeds	Configured; awaiting Colab budget (must confirm v2 does not re-introduce divergence)	—
D2. Parity-tuned head-to-head vs FedKSeed ($lr \times \beta$ grid, 1000 rounds)	Default-hyperparameter run done (§6.5); tuned run pending	—
D3. Combo + DP composition: D-MeZO-N v2 + $\epsilon=10$ via RDP accountant	Configured; awaiting Colab budget	—

E. Honest framing for the strongest publishable claim

Based on solidly-demonstrated tier A + theoretical contribution:

D-MeZO-N v2 = combo (adaptive p-clip B1 + drift-reset B5) is the first decentralized federated zeroth-order optimizer for LLM fine-tuning with (i) a paper-scale multi-seed validated loss improvement over vanilla MeZO (Qwen3.5-4B-Base / MathLogicQA / 3 paired seeds: $\Delta\text{loss} = -5.5\%$, 3/3 same direction, paired $t = 5.3$, $p < 0.05$; lowest loss-std among the variants that improve on vanilla), (ii) closed-form Lyapunov stability proofs under PL + heavy-ball + β -decay + p-clipping (Theorem 3 — the first such guarantee in the ZO-LLM

setting), (iii) a formal per-round ($\varepsilon = 10, \delta = 10^{-3}$)-DP guarantee via natural ρ -clip sensitivity bound at $\sim 6\%$ utility cost (Theorem 4 + §6.7), (iv) $\approx 5 \times 10^8 \times$ (roughly 8–9 orders of magnitude: 8 GB vs 16 B per round) communication compression vs. FedAvg (1 scalar + 1 seed per round) plus independent per-client z_i giving $1/n$ variance reduction on both noise components, and (v) demonstrated convergence on hybrid linear-attention LLMs across complete / ring / non-IID-Dir(0.5) topologies, beating FedKSeed on loss 3/3 seeds (§6.5).

What is **not** claimed (per tier C): asymptotic acceleration over vanilla MeZO, $O(1/T^2)$ rates, multi-direction strict improvement, accuracy gains beyond seed noise on convergent tasks (the +2.3pp accuracy trend is reported for transparency only, paired $t \approx 0.8, p \approx 0.5$). What is **not yet ready** (tier B/D): rescue-regime safety of the v2 combo on HellaSwag, a parity-tuned FedKSeed comparison, and T -round DP composition via an RDP accountant.

9. Conclusion

We presented D-MeZO-N — Decentralized Federated MeZO with Nesterov-style **stabilization** — and established it as a viable peer-to-peer federated optimizer for LLM fine-tuning. Its eight contributions (C1–C8) cover novel hybrid-linear-attention architecture support, robustness to extreme non-IID, negligible topology cost, the **D-MeZO-N v2 (combo B1+B5)** recipe with a paper-scale multi-seed validated **loss** improvement over vanilla MeZO ($\Delta_{\text{loss}} = -5.5\%$, 3/3 seeds, paired $t = 5.3, p < 0.05$) and a 3/3-seed loss win over FedKSeed-at-default, four formal convergence theorems (T1–T4), and the first formal per-round ($\varepsilon = 10, \delta = 10^{-3}$)-DP guarantee for decentralized federated ZO on LLMs. Theorem 3 is, to our knowledge, the first stability guarantee for ZO heavy-ball in the PL regime; its contraction rate matches plain SGD-under-PL — momentum is made *safe*, not *faster*, in the high-variance ZO regime. The accuracy trend (+2.3pp) is reported for transparency only and not claimed as a gain. Open directions: scaling to 100+ clients, multi-seed validation of the rescue regime, generative tasks, a parity-tuned FedKSeed comparison, and T -round DP composition via an RDP accountant with subsampling amplification.

Code and data availability. All code, configs, tests and the committed raw result JSONs are available in the companion repository github.com/Siesher/dmezo. File paths of the form `docs/...`, `scripts/...`, `tests/...`, `notebooks/...` referenced throughout the paper point to files in that repository; in particular, the consolidated proofs are in `docs/theory_rigorous.md` and the headline raw data in `experiments/diagnostics/validate_multiseed_fed_Qwen_Qwen3p5-4B-Base_mathlogicqa.json` and `head-to-head/head_to_head_fedkseed_Qwen_Qwen3p5-4B-Base_mathlogicqa.json`.

References

- Abadi, M., Chu, A., Goodfellow, I., McMahan, H. B., Mironov, I., Talwar, K., Zhang, L. (2016). Deep learning with differential privacy. ACM CCS 2016. arXiv:1607.00133.
- Aybat, N. S., Fallah, A., Gurbuzbalaban, M., Ozdaglar, A. (2019). A universally optimal multistage accelerated stochastic gradient method. NeurIPS 2019.
- Beznosikov, A., Richtárik, P., Diskin, M., Ryabinin, M., Gasnikov, A. (2022). Distributed methods with compressed communication for solving variational inequalities, with theoretical guarantees. NeurIPS 2022. arXiv:2207.10792.
- Bottou, L., Curtis, F. E., Nocedal, J. (2018). Optimization methods for large-scale machine learning. SIAM Review 60(2):223–311. arXiv:1606.04838.

- Dwork, C., Roth, A. (2014). The algorithmic foundations of differential privacy. *Foundations and Trends in Theoretical Computer Science* 9(3–4):211–407.
- Dwork, C., Rothblum, G. N., Vadhan, S. (2010). Boosting and differential privacy. *FOCS 2010*.
- Gasnikov, A., Dvinskikh, D., Dvurechensky, P., Gorbunov, E., Beznosikov, A., Lobanov, A. (2023). Randomized gradient-free methods in convex optimization. *Encyclopedia of Optimization* (Springer), pp. 1–15. arXiv:2211.13566.
- Karimi, H., Nutini, J., Schmidt, M. (2016). Linear convergence of gradient and proximal-gradient methods under the Polyak-Łojasiewicz condition. *ECML-PKDD 2016*.
- Koloskova, A., Loizou, N., Boreiri, S., Jaggi, M., Stich, S. U. (2020). A unified theory of decentralized SGD with changing topology and local updates. *ICML 2020*.
- Li, Z., Ying, B., Liu, Z., Dong, C., Yang, H. (2024). Achieving dimension-free communication in federated learning via zeroth-order optimization. *ICLR 2025*. arXiv:2405.15861.
- Liu, C., Zhu, L., Belkin, M. (2022). Loss landscapes and optimization in over-parameterized non-linear systems and neural networks. *Applied and Computational Harmonic Analysis* 59:85–116.
- Malladi, S., Gao, T., Nichani, E., Damian, A., Lee, J. D., Chen, D., Arora, S. (2023). Fine-tuning language models with just forward passes. *NeurIPS 2023*.
- Maritan, A., Ridolfi, A., Notarstefano, G. (2024). FedZeN: a zeroth-order Newton-style method for federated learning. arXiv:2309.17241.
- Mironov, I. (2017). Rényi differential privacy. *IEEE Computer Security Foundations Symposium (CSF) 2017*. arXiv:1702.07476.
- Nesterov, Y., Spokoiny, V. (2017). Random gradient-free minimization of convex functions. *Foundations of Computational Mathematics* 17(2):527–566.
- Polyak, B. T. (1964). Some methods of speeding up the convergence of iteration methods. *USSR Computational Mathematics and Mathematical Physics* 4(5):1–17.
- Qin, Z., Chen, D., Qian, B., Ding, B., Li, Y., Deng, S. (2024). Federated full-parameter tuning of billion-sized language models with communication cost under 18 kilobytes. *ICML 2024*. arXiv:2312.06353.
- Sadegh, P., Spall, J. C. (1998). Optimal random perturbations for stochastic approximation using a simultaneous perturbation gradient approximation. *IEEE Transactions on Automatic Control* 43(10):1480–1484.
- Sadiev, A., Borodich, E., Beznosikov, A., Dvinskikh, D., Chezhegov, S., Tappenden, R., Takáč, M., Gasnikov, A. (2022). Decentralized personalized federated learning: Lower bounds and optimal algorithm for all personalization modes. *EURO Journal on Computational Optimization* 10:100041.
- Shu, Y., Yao, W., Hu, S. X. (2024). Ferret: federated full-parameter tuning at scale for large language models. arXiv:2409.06277.
- Spall, J. C. (1992). Multivariate stochastic approximation using a simultaneous perturbation gradient approximation. *IEEE TAC* 37(3):332–341.

- Spall, J. C. (1997). A one-measurement form of simultaneous perturbation stochastic approximation. *Automatica* 33(1):109–112.
- Spall, J. C. (2000). Adaptive stochastic approximation by the simultaneous perturbation method. *IEEE Transactions on Automatic Control* 45(10):1839–1853.
- Stich, S. U. (2019). Local SGD converges fast and communicates little. *ICLR* 2019.
- Yang, T., Lin, Q., Li, Z. (2016). Unified convergence analysis of stochastic momentum methods. *arXiv:1604.03257*.
- Zellers, R., Holtzman, A., Bisk, Y., Farhadi, A., Choi, Y. (2019). HellaSwag: Can a machine really finish your sentence? *ACL* 2019. *arXiv:1905.07830*.